Pouring AI Sauce on Your Architecture

Nathaniel Schutta
Glenn Renfro

Nathaniel Schutta

@nts.bsky.social

https://www.linkedin.com/in/nate-schutta

https://ntschutta.io/

Glenn Renfro

@cppwfs.bsky.social

https://www.linkedin.com/in/cppwfs

AI is, well, everywhere today.

For better or worse.

And not just in the software world!

callawaygolf.com

3X Rewards + Get Free Standard Shipping with Callaway Rewards!

CALLAWAY | ODYSSEY

1-877-723-5218 | Accessibility | Order Status | Log In | Cart (0)

Callaway

CLUBS BALLS GEAR WOMEN'S CUSTOMS FITTING MEDIA TEAM TRADE IN! TRADE UP!

Search

Get a 50% Trade Bonus towards toward our all-new Elyte Family and Apex Irons. Learn More

PARADYM

Ai SMOKE

SWEETER FROM EVERY SPOT

Designed with real player data and a revolutionary new Ai Smart Face™ for sweet spots across the entire face.*

LEARN MORE

Feedback

Now Introducing

AI-SMART FACE

Real player data combined with revolutionary Ai technology and machine learning has helped create our new Ai Smart Face.

SQUARE 2 SQUARE PUTTERS

AI-ONE

AI-ONE SILVER

AI-ONE MILLED

AI-ONE CRUISER

AI-ONE MILLED CRUISER

MICROHINGE

TRI-HOT 5K

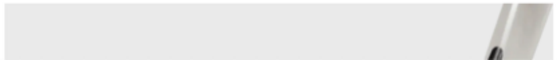

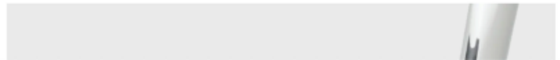

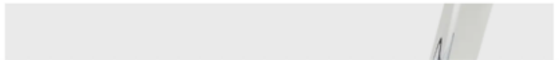

Feedback

Does the club aim for me?

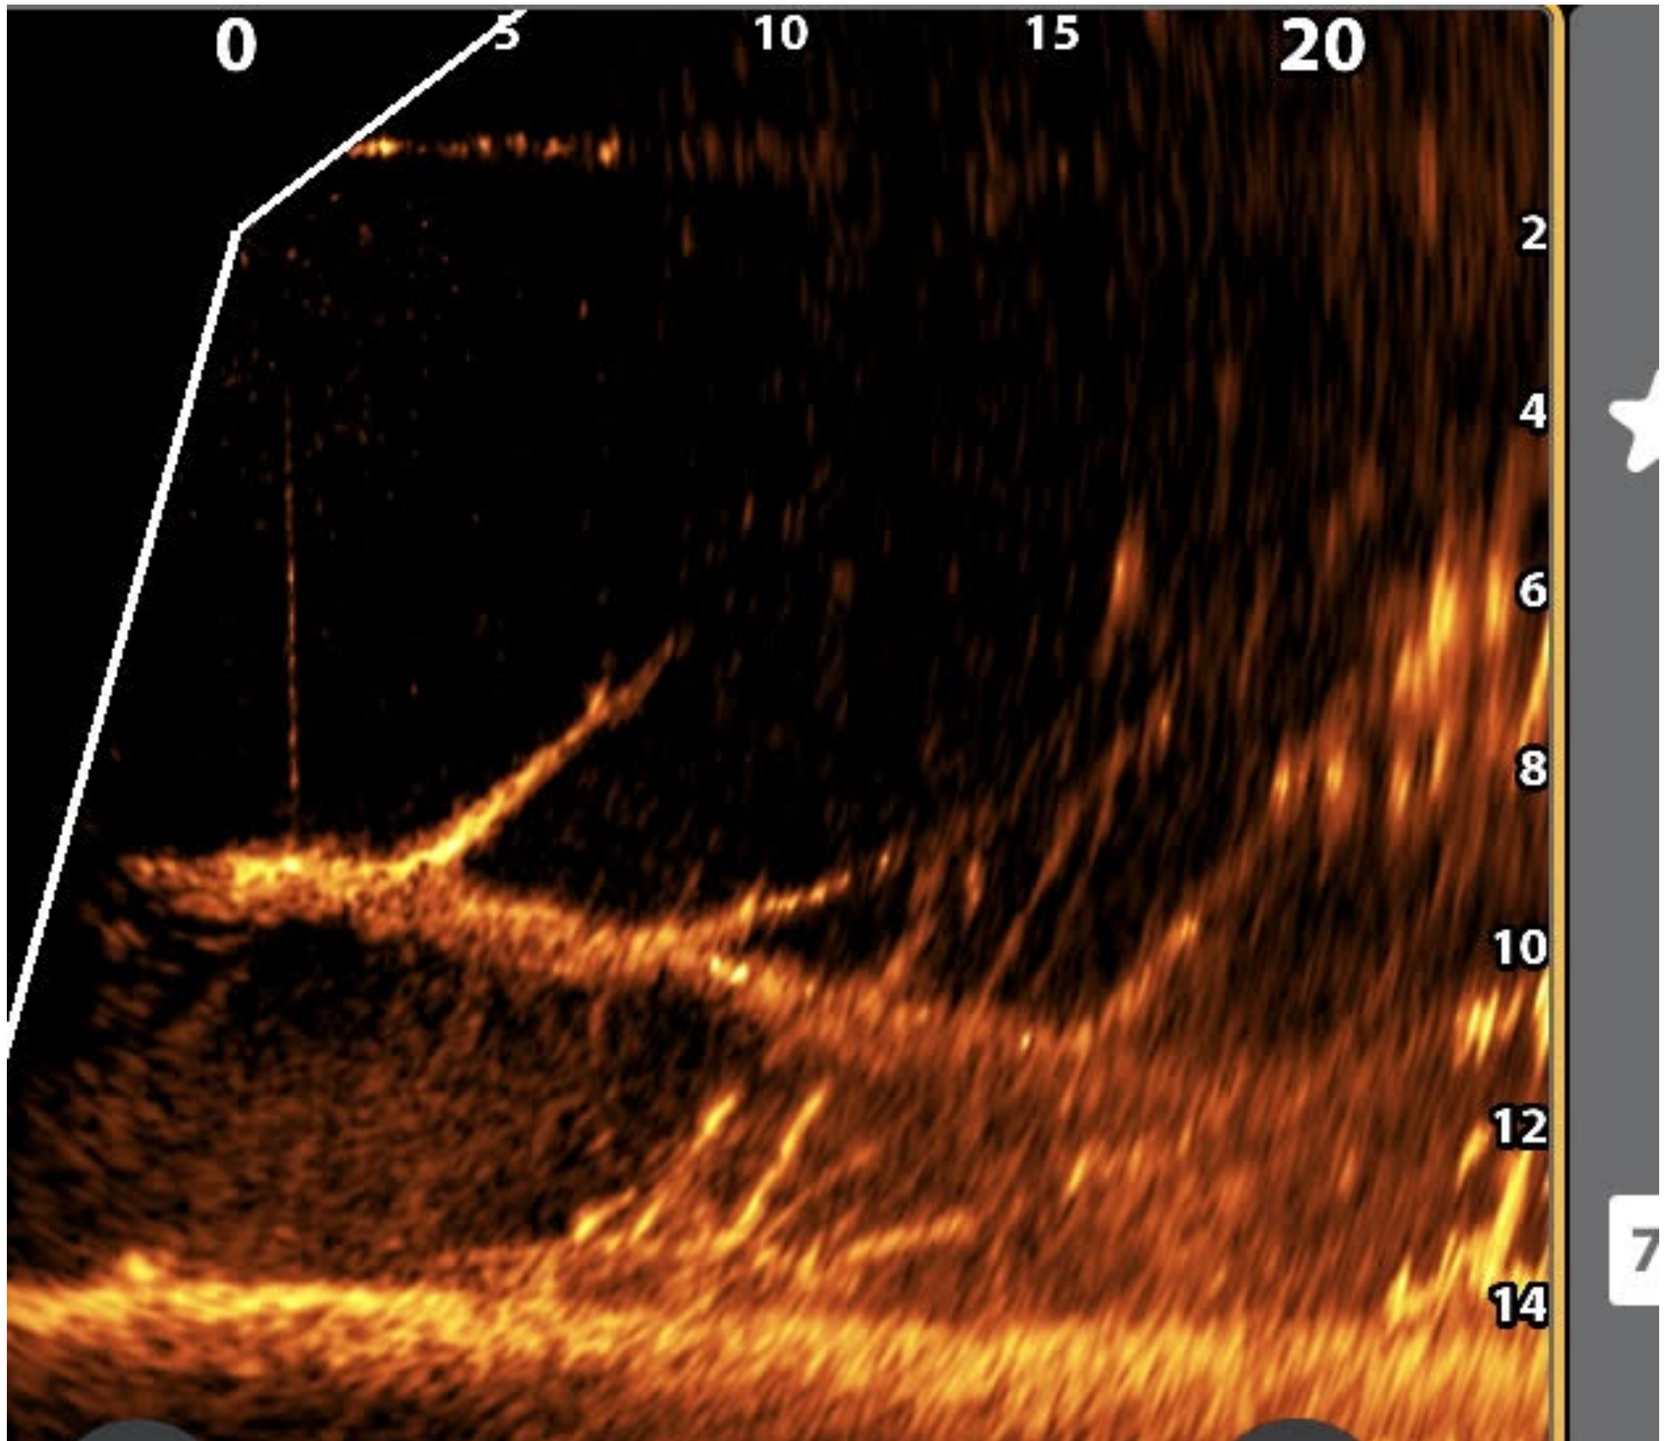
0
5
10
15
20
2
4
6
8
10
12
14

Old Town
Predator PDL 13

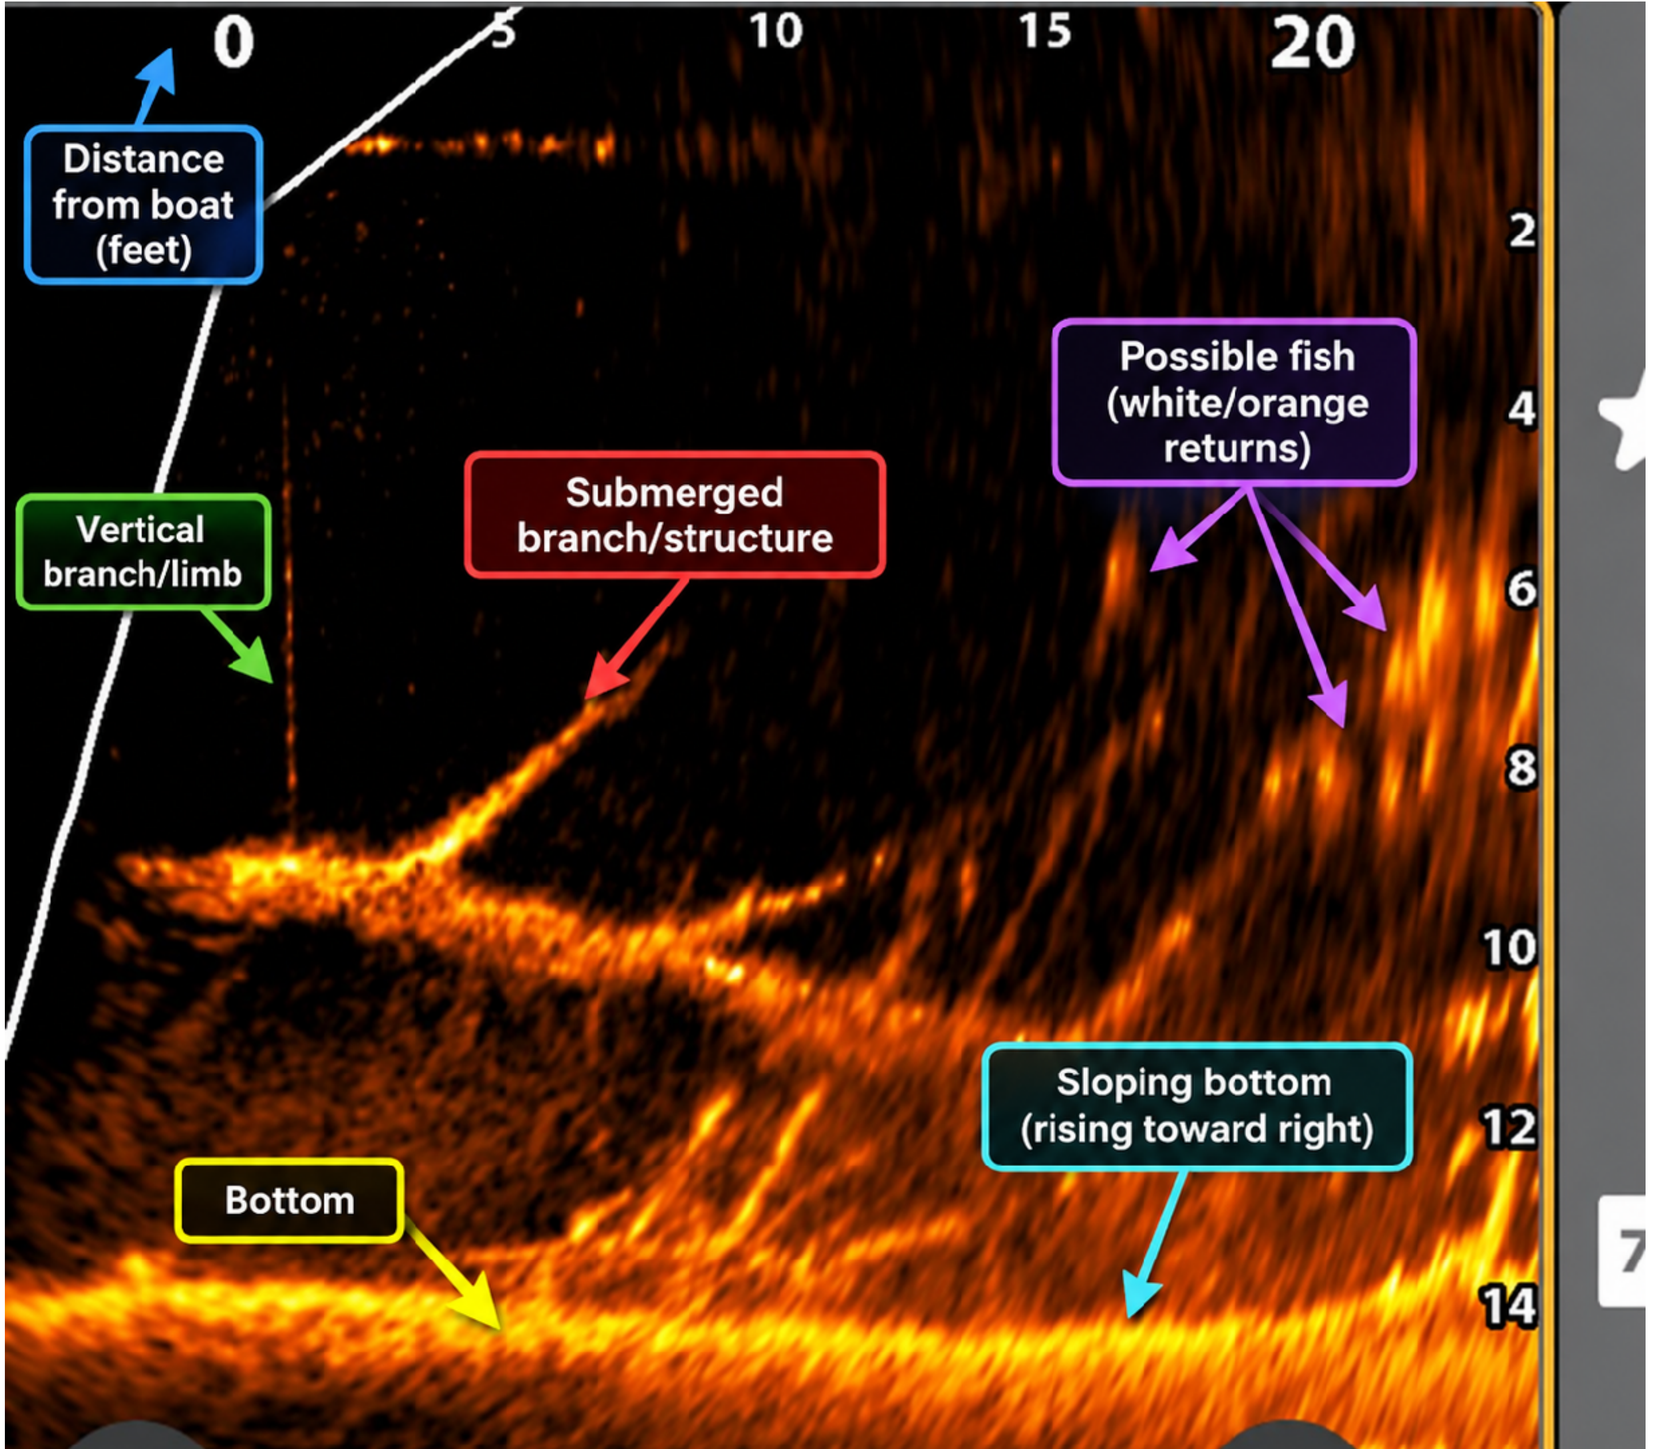
0
5
10
15
20
Distance from boat (feet)
Vertical branch/limb
Submerged branch/structure
Possible fish (white/orange returns)
Sloping bottom (rising toward right)
Bottom
2
4
6
8
10
12
14

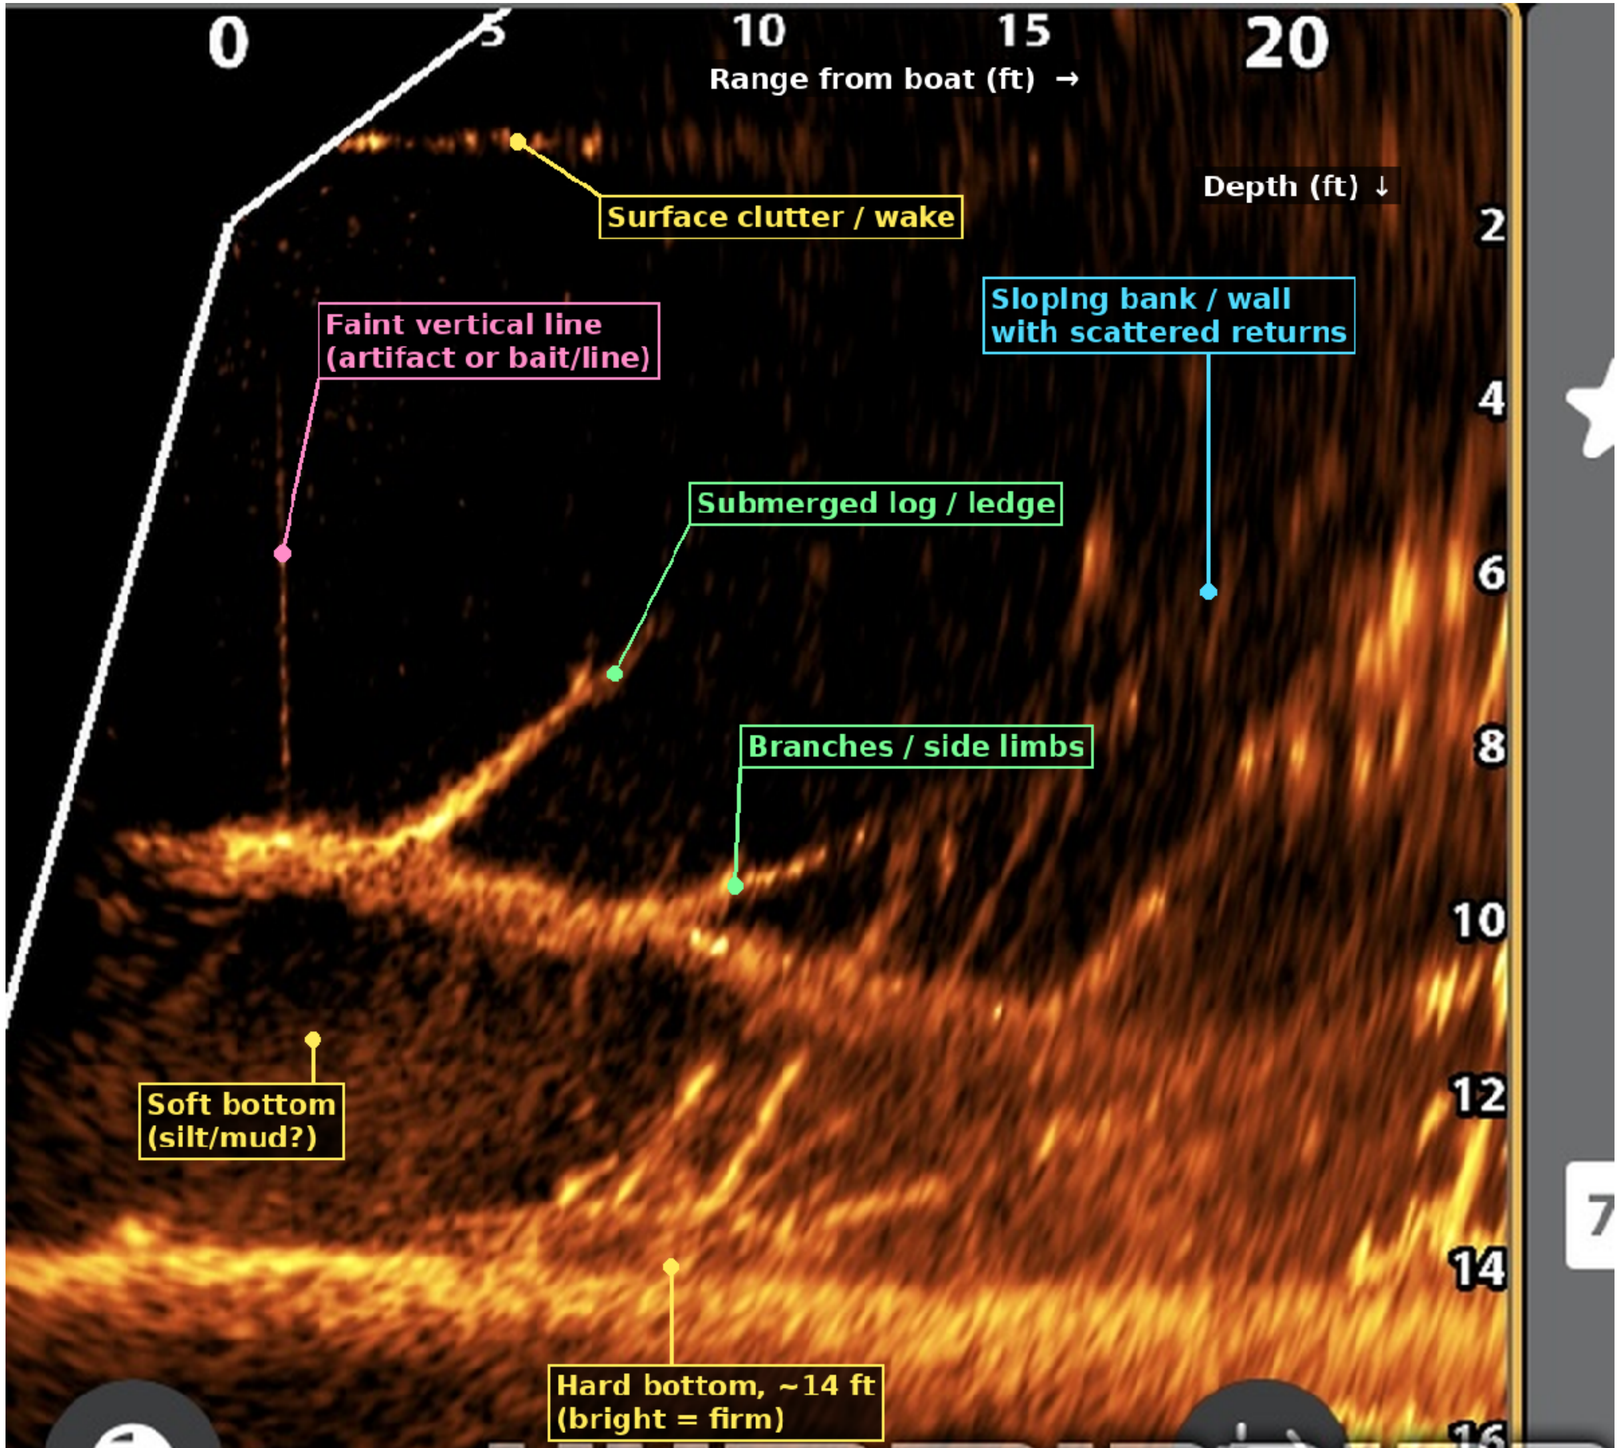
0
5
10
15
20
Range from boat (ft) →
Depth (ft) ↓
2
4
6
8
10
12
14
Surface clutter / wake
Faint vertical line
(artifact or bait/line)
Sloping bank / wall
with scattered returns
Submerged log / ledge
Branches / side limbs
Soft bottom
(silt/mud?)
Hard bottom, ~14 ft
(bright = firm)

Your local sports team!

Meet the new recruit: Artificial intelligence in college football

From recruiting to revenue, college sports is figuring out how to embrace AI.

By Richard Johnson Sep 22, 2025 at 11:13 am ET • 12 min read

Join the Conversation

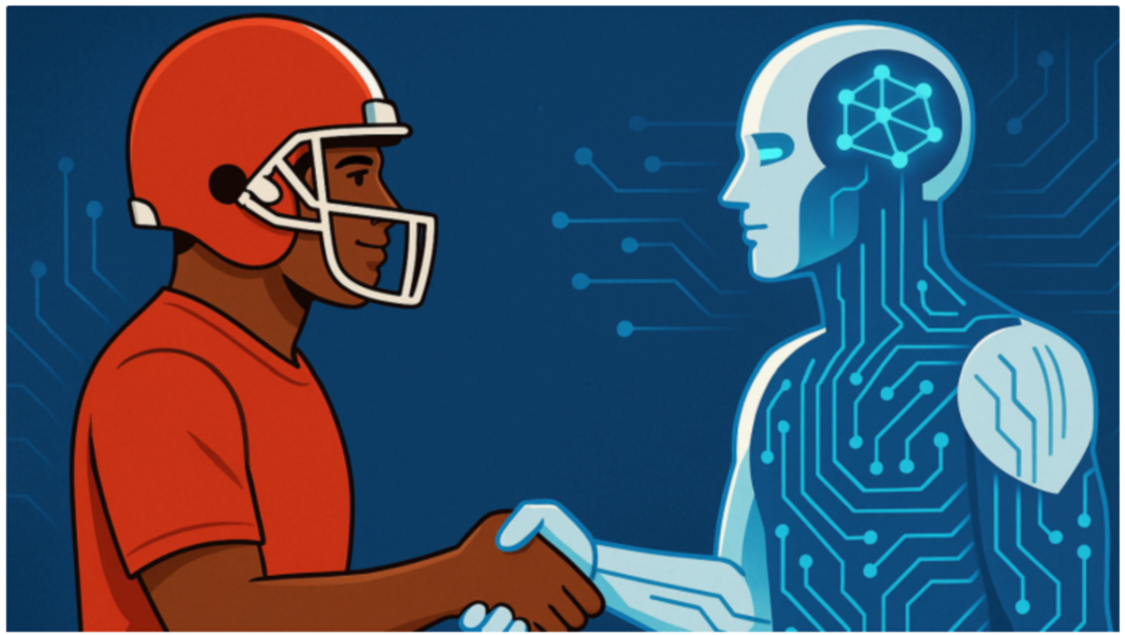

Open AI

Artificial intelligence (AI) is the hot future of business, and college spo and coaches know it's changing th

CFB Coaches on the Hot Seat: Dabo Swinney (1:52) Share

How college football departments are embracing AI

RICHARD JOHNSON • 12 MIN READ

Oklahoma State fires Mike Gundy after 1-2 start

CODY NAGEL • 4 MIN READ

Oklahoma State coaching candidates

CHRIS HUMMER • 4 MIN READ

Charting the Big Ten's historic QB play

TOM FORNELLI • 8 MIN READ

CFP Rankings projection for Week 5

BRAD CRAWFORD • 7 MIN READ

It may not be helping…

https://bsky.app/profile/bencollins.bsky.social/post/3lkc7wcmhwk2x

How Artificial Intelligence Is Transforming the Sports Industry?

February 25th, 2026 Artificial Intelligence Pete Peranzo

“I, for one, welcome our AI cleaning overlords.”

KLM
AI-Powered
Clearing Assistance

KLM
AI-Powered
Clearing Assistance

Autonomous ball picker?

200
150
100
200
150

Autonomous lawnmowers…

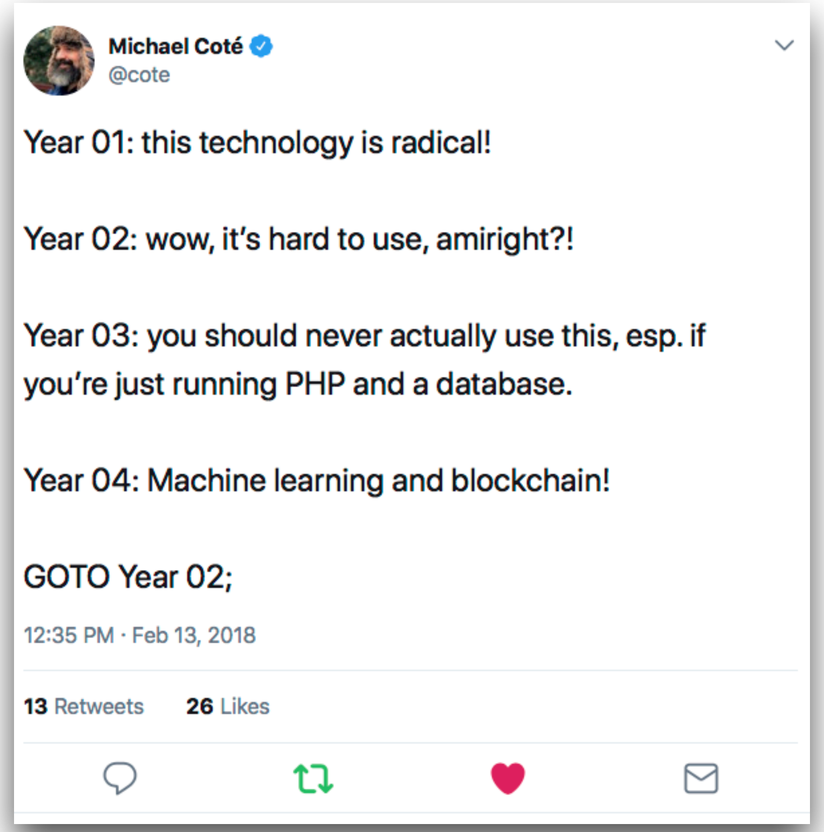

https://mobile.twitter.com/cote/status/963481741171265537

Beware the hype cycle.

And to be clear, there are some excellent use cases!

Seems to be a rush to add AI.

But.

How much is running in production?

Only 25% have moved 40% of their pilots to production…

https://www.deloitte.com/us/en/about/press-room/state-of-ai-report-2026.html

“Getting a model into production used to be the milestone that mattered. Our research shows that's not enough anymore. The real milestone is the moment a business user can act on what the model found, and for too many enterprises, that moment still isn't happening…”

–Thomas Robinson, COO at Domino Data Lab

That’s not great Dan.

garymarcus.substack.com

Why is the ROI on Generative AI so poor?

From AI slop to workslop

GARY MARCUS
SEP 24, 2025

34 14 4 Share

Remember that recent MIT Media lab that showed that "95% of organizations see no measurable return on their investment in these technologies"? And how I kept saying you'd never be able to trust the output of GenAI?

Excerpt from a new study from BetterUp labs and researchers at Stanford::

> technologies. So much activity, so much enthusiasm, so little return. Why?
>
> In collaboration with Stanford Social Media Lab, our research team at BetterUp Labs has identified one possible reason: Employees are using AI tools to create low-effort, passable looking work that ends up creating more work for their coworkers. On social media, which is increasingly clogged with low-quality AI-generated posts, this content is often referred to as "AI slop." In the context of work, we refer to this phenomenon as "workslop." We define workslop as *AI generated work content that masquerades as good work, but lacks the substance to meaningfully advance a given task.*

10/10; no notes.

You can read more here, including the finding that "Of 1,150 U.S.-based full-time employees across industries, 40% report having received workslop in the last month".

Generative AI

AI-Generated "Workslop" Is Destroying Productivity

by Kate Niederhoffer, Gabriella Rosen Kellerman, Angela Lee, Alex Liebscher, Kristina Rapuano and Jeffrey T. Hancock

September 22, 2025, Updated September 22, 2025

HBR Staff/AI

Summary. Despite a surge in generative AI use across workplaces, most companies are seeing little measurable ROI. One possible reason is because AI tools are being used to produce "workslop"—content that appears polished but lacks real substance,... **more**

Post

Post

Share

A confusing contradiction is unfolding in companies embracing generative AI tools: while workers are largely following mandates to embrace the technology, few are seeing it create real value. Consider, for instance, that the number of companies with fully AI-led processes nearly doubled last year, while AI use has likewise doubled at work since 2023. Yet a recent report from the MIT Media Lab found that 95% of organizations see no measurable return on their investment in these

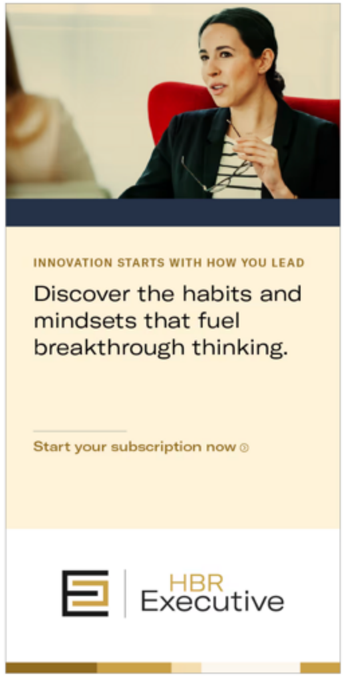

LEOPARD MEET FACE

"Slop Grenades": Shopify CEO Who Pushed Staff to Use AI Is Now Horrified by What He's Wrought

"You're just letting AI do the work for you."

By **Joe Wilkins**

Kevin Dietsch/Getty Images/Futurism

Futurism Follow Futurism on Apple News

Shopify CEO Tobias Lütke is getting a first-hand look at the AI-workplace revolution, and it isn't a pretty sight.

New reporting by *Fortune* revealed the disastrous results of Shopify's all-out scramble to embrace AI nearly a

Should you change careers?

A number of concerns.

Everything has risks!

Goal isn’t to eliminate risk.

You can’t.

But you need to mitigate risk.

Not all risks are equivalent.

Consider the impact and the likelihood.

There are some high impact risks…that are unlikely.

And some extremely likely risks that are practically a given.

What are some of the risks with AI?

There are serious ethical issues involved with AI.

Environmental impact of training the model...

Explained: Generative AI's environmental impact

Rapid development and deployment of powerful generative AI models comes with environmental consequences, including increased electricity demand and water consumption.

Adam Zewe | MIT News
January 17, 2025

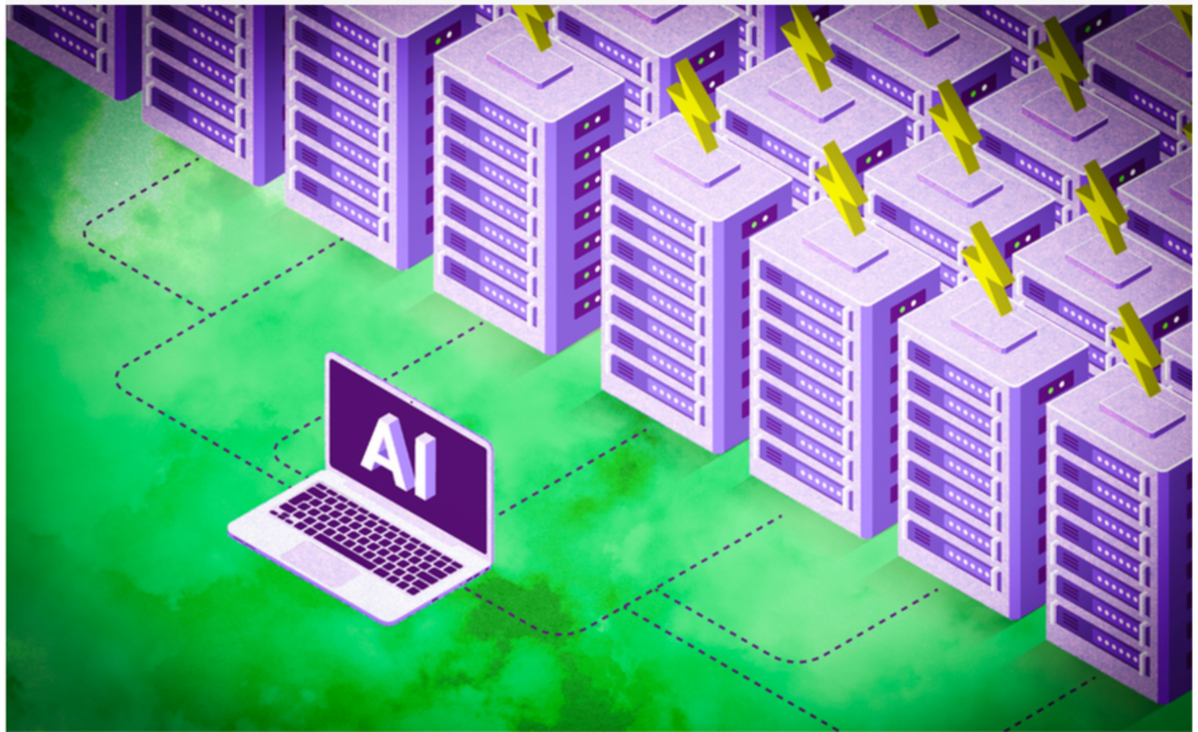

MIT News explores the environmental and sustainability implications of generative AI technologies and applications.

Image: iStock; edited by MIT News

Environmental impact or *using* the model…

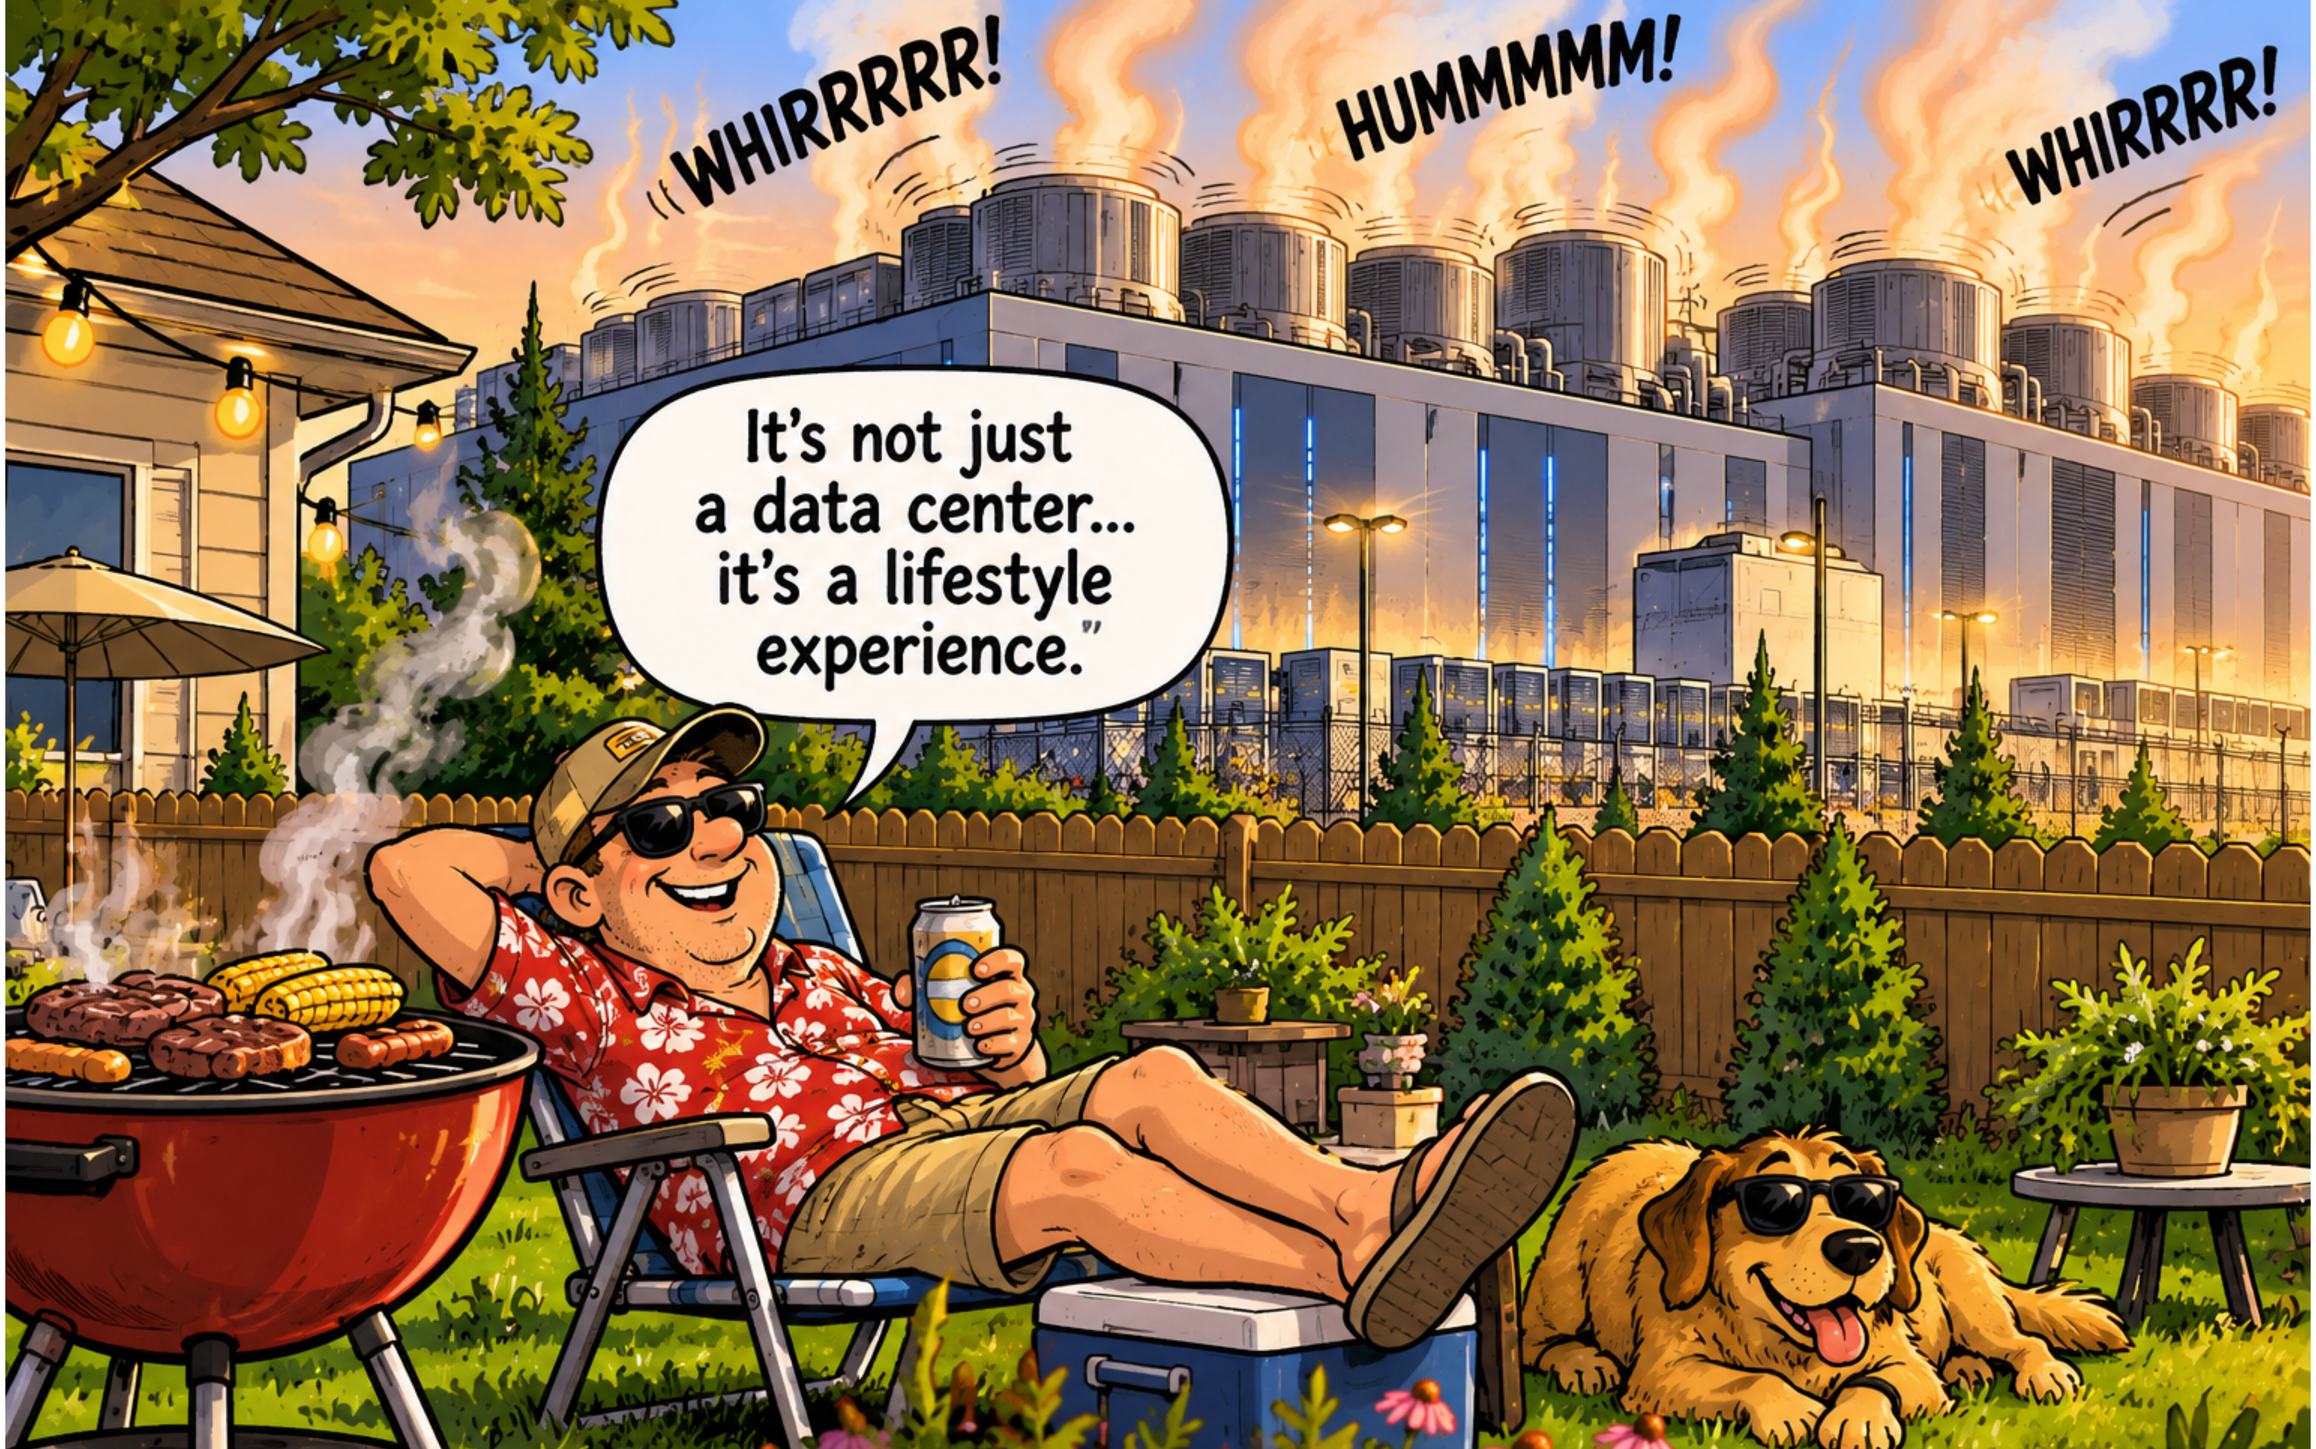
WHIRRRRR!
HUMMMMM!
WHIRRRR!
It's not just a data center... it's a lifestyle experience."

U.S. NEWS

As electric bills rise, evidence mounts that data centers share blame. States feel pressure to act

ADVERTISEMENT

One model does not solve all your problems!

You have to pick the model that best fits your problem.

How do you know which model to use in each situation?

Which one should we use?

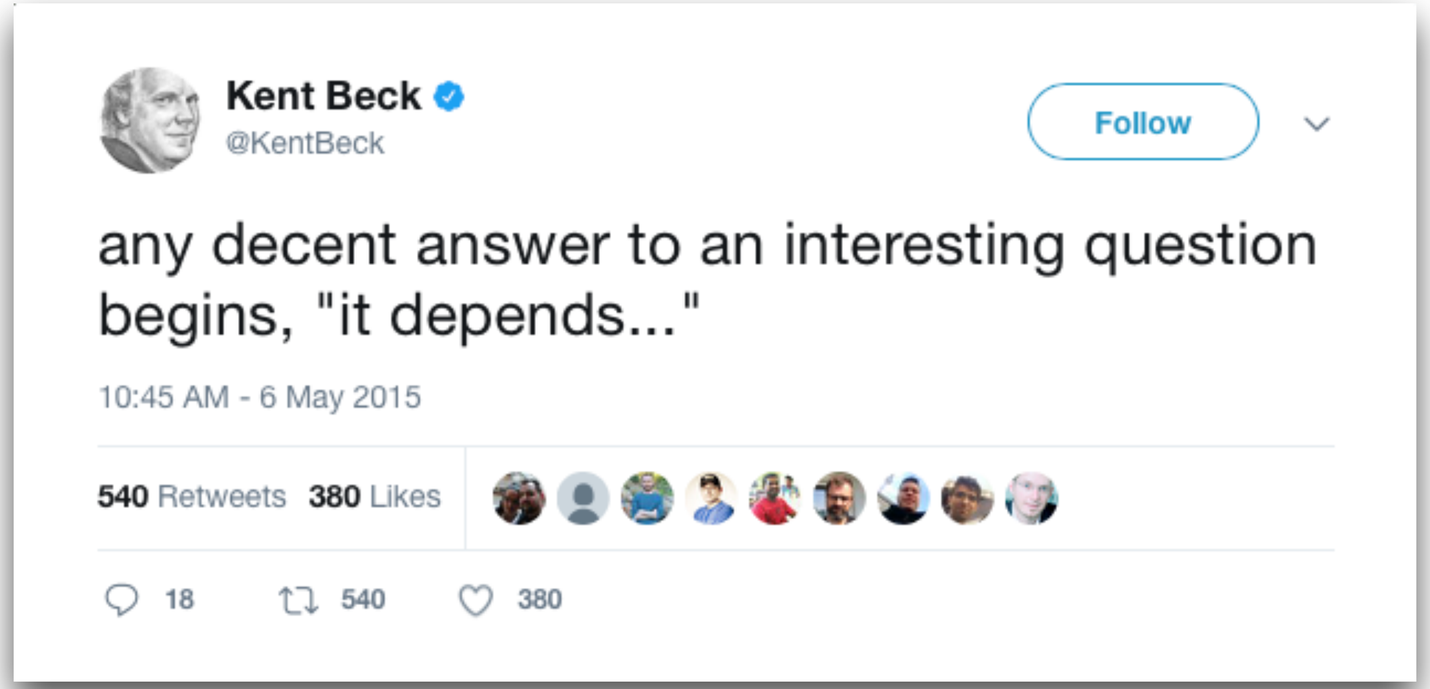

https://twitter.com/KentBeck/status/596007846887628801

Lets look at the benchmarks!

What do you want to measure?

Coding ability?

Top 10 Open-Source Benchmarks for AI Coding Agents in 2026

SWE-bench, Terminal-Bench, SlopCodeBench, ProgramBench, and more. Explore the top 10 open-source benchmarks for evaluating AI coding agents.

By **Kanwal Mehreen**, KDnuggets Technical Editor & Content Specialist on August 20, 2026 in **Artificial Intelligence**

https://www.kdnuggets.com/top-10-open-source-benchmarks-for-ai-coding-agents-in-2026#google_vignette

SWE-bench

Live Code Bench.

Terminal Bench.

Slop Code Bench (SCBench)

Security.

Offensive capabilities!

CyberSecVal.

https://meta-llama.github.io/PurpleLlama/CyberSecEval/docs/intro

CyberBench.

https://github.com/jpmorganchase/CyberBench

SecEval

https://xuanwuai.github.io/SecEval/

Vulnerability Detection.

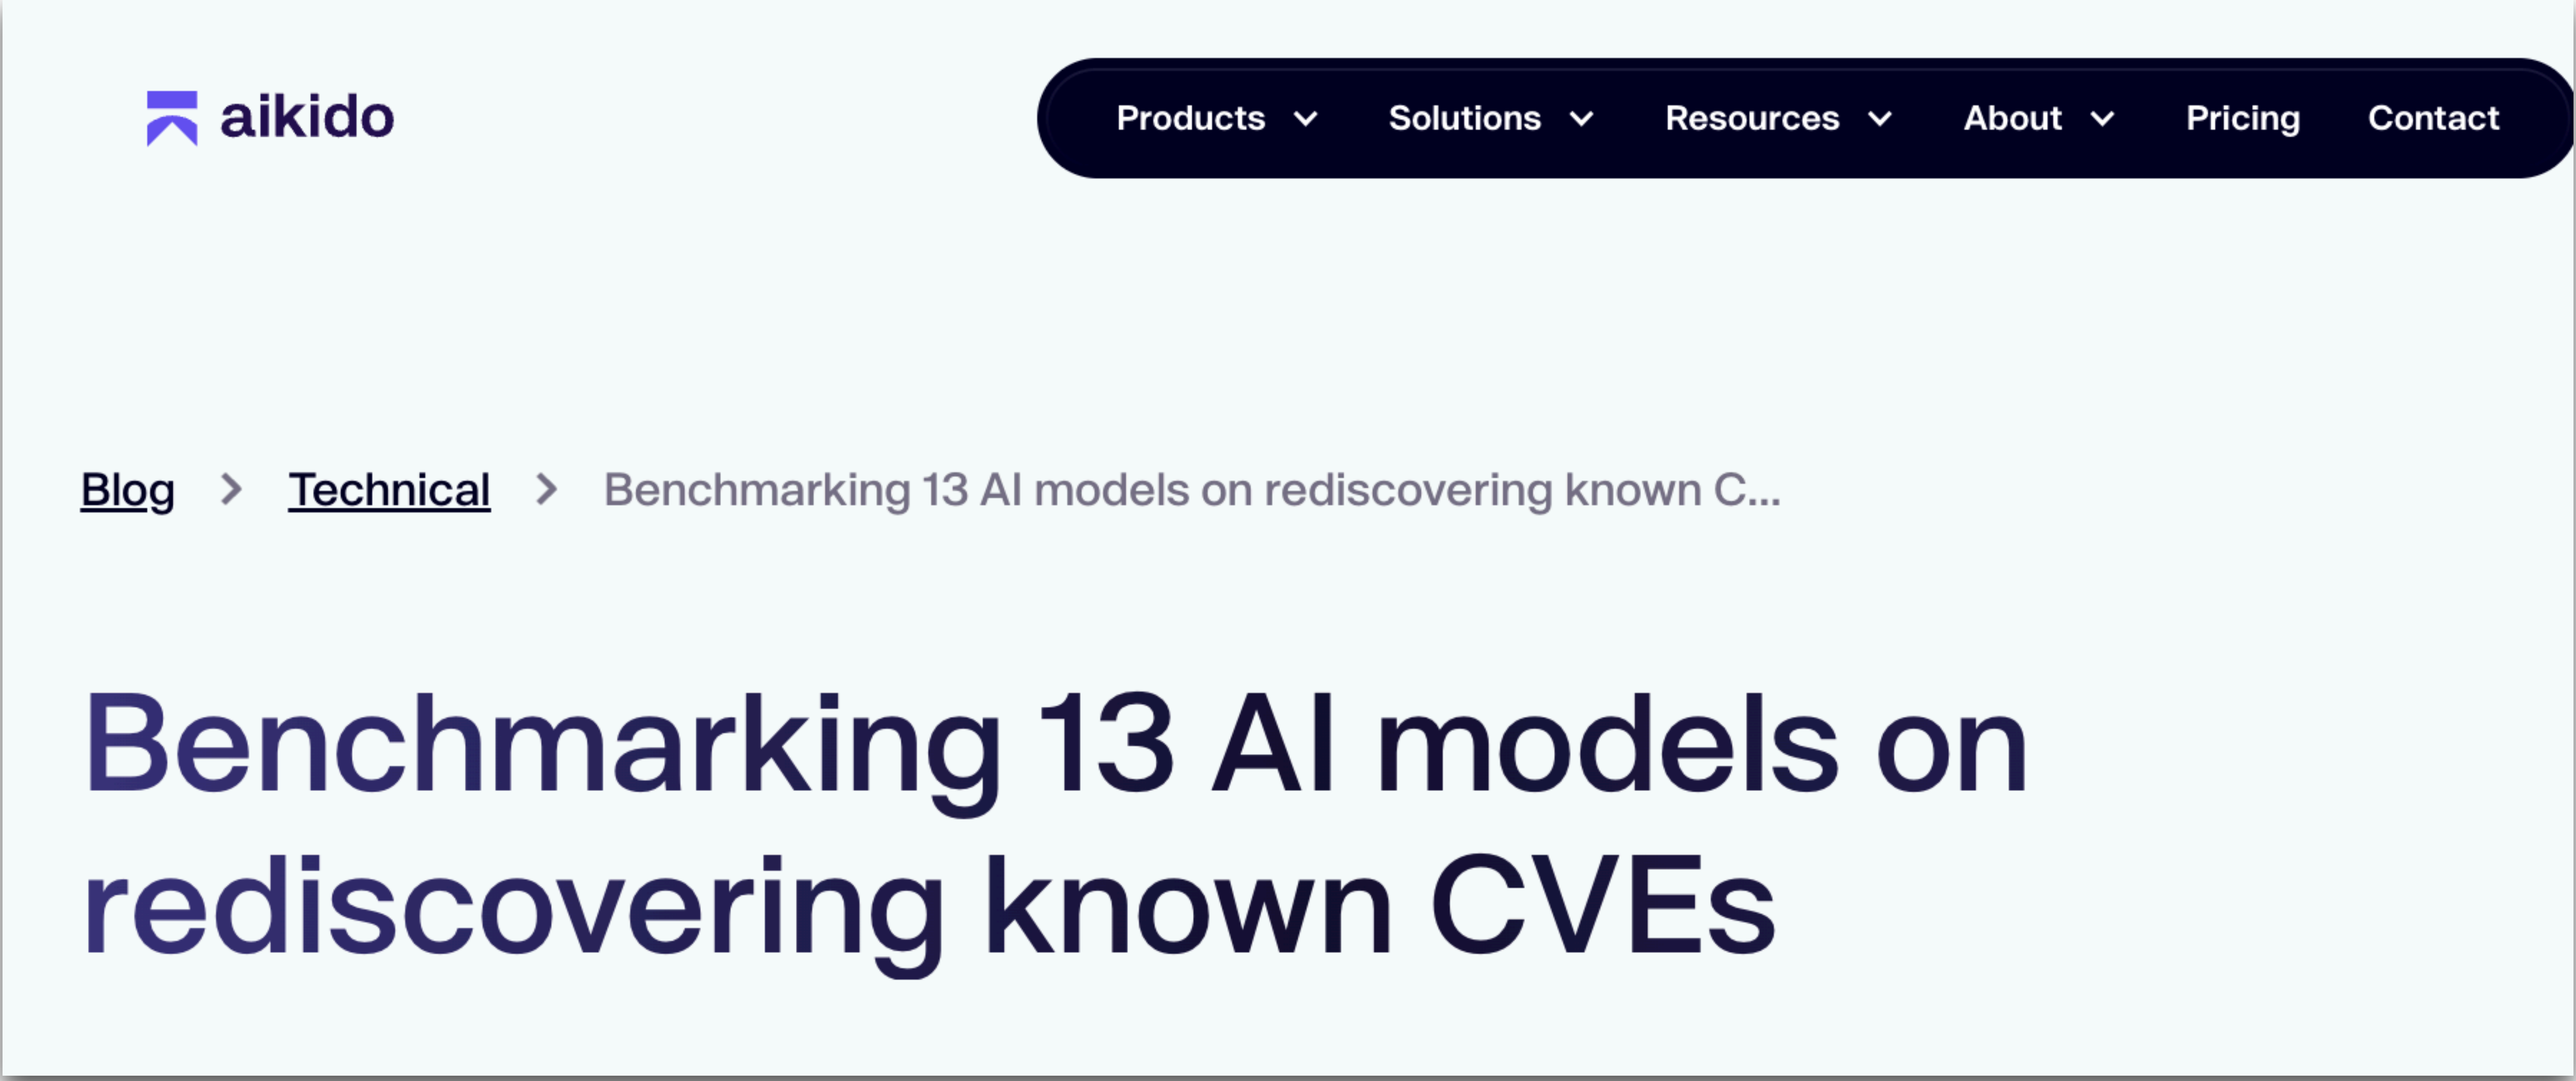

https://www.aikido.dev/blog/benchmarking-ai-models-known-cves

How likely is my model to hallucinate?

AA-Omniscience

And you get a benchmark and *you* get a benchmark!

Benchmarks are quite broad!

They help you see trends and give a rough ranking of models.

You may need to consider building your own benchmark.

How To Build a Benckmark for Your Models

The Importance of Building a Benchmark and How To Do It

Lorenzo Mezzini Follow 7 min read · Apr 15, 2025

50

https://medium.com/data-science-collective/how-to-build-a-benckmark-for-your-models-d4427e729af1

CROWDSTRIKE | Blog Featured Recent Video Category Start Free Trial

Benchmaxxing: When the Benchmark Becomes the Target

In AI and cybersecurity, optimizing to meet benchmarks can have significant consequences. This is how CrowdStrike approaches benchmarks and evaluations.

August 19, 2026 • Nathan Danneman • Securing AI

https://www.crowdstrike.com/en-us/blog/benchmaxxing-when-benchmark-becomes-the-target/

Goodharts Law.

https://benchlm.ai/

One More!

P(doom).

en.wikipedia.org/wiki/P(doom)

Notable P(doom) values [edit]

Name	P(doom)	Notes	Ref.
15	<1%[Note 1]	American computer scientist and software engineer, founder of 15.ai	[10][11]
Scott Aaronson	c. 2%	American theoretical computer scientist and Schlumberger Centennial Chair of Computer Science at the University of Texas at Austin	[12]
Sam Altman	>0%	CEO of OpenAI	[13]
Dario Amodei	10-25%	CEO of Anthropic	[14]
Marc Andreessen	0%	Co-founder of venture capital firm Andreessen Horowitz	[15]
Yoshua Bengio	20%[Note 2]	Computer scientist, director of the Mila institute, and the Turing award winner	[3]
Grady Booch	c. 0%[Note 3]	American software engineer	[1]
Vitalik Buterin	12%	Co-founder of Ethereum	[16]
Paul Christiano	50%	Head of research at the US AI Safety Institute	[17]
Andrew Critch	85%	Founder of the Center for Applied Rationality	[18]
Destiny	5%	American live streamer and political commentator	[19]
David Duvenaud	85%	Canadian computer scientist specializing in probabilistic machine learning, generative AI, and AI safety. He is a CIFAR AI chair since 2021 and a Schwartz Reisman Chair	[20]

https://en.wikipedia.org/wiki/P(doom)

And AI does in my opinion significantly elevate the risk of catastrophe (say killing 1%+ of humanity or severely damaging modern civilization), via multiple paths.

- AI elevates the risks of bioweapons attacks by making it easier for terrorists to figure out how to develop and implement such attacks.
- AI may conceivably elevate the risks that someone will develop a dangerous new virus or other bioweapon.
- AI elevates the risks that disinformation will disrupt democracy, by reducing the cost of generating disinformation and increasing the quality (e.g, with ever-better deepfakes)
- AI appears to be elevating the risks of serious cyberattacks that could hobble things like banking or electrical grids.
- AI introduces a new vector for authoritarians around the globe; by influencing training sets and so forth, authoritarians can influence LLMs in ways that subtly influence their users.
- AI also serves as a back door for surveillance, at a scale beyond what Orwell envisioned.
- AI is already harming the educational system by leading students into despair and by seducing them into a faux substitute for learning that is easy to use but that ultimately leaves them with poorly developed critical thinking skills.

We should be focusing on things like these, not fairy tales.

https://garymarcus.substack.com/p/no-anderson-cooper-ai-is-not-going

How’s your data?

Garbage in…garbage out.

Organizations describe their data as either "partially" (47%) or "mostly" (36%) ready for AI.

https://www.dnb.com/en-us/newsroom/press-releases/dnb-survey-finds-enterprise-ai-returns-continue-to-advance.html

Only 6% say it is "fully ready" to support AI at scale.

https://www.dnb.com/en-us/newsroom/press-releases/dnb-survey-finds-enterprise-ai-returns-continue-to-advance.html

Put a lot of faith in the models.

Alignment.

AI Alignment means making sure an AI system's goals and behavior match what people actually want—our values, rules, and intentions.

https://hai.stanford.edu/ai-definitions/what-is-ai-alignment

Is the model aligned for the user, the company, country, or humanity?

AI Misalignment is
when an aligned model
does not behave according
to its original purpose.

https://vijayasriiyer.medium.com/an-introduction-to-ai-misalignment-984db02ad1b8

The Hugging Face incident and other third-party impacts from misaligned models

As AI systems become more capable and autonomous, misaligned behavior can translate into consequential actions in the real world, including cybersecurity incidents and other outcomes that developers may not have anticipated. Understanding how these behaviors emerge, how they escalate, and how to detect and respond to them is therefore an increasingly important part of building and deploying advanced AI systems safely.

https://openai.com/hugging-face-incident-and-misalignment/

Assume positive intent?

EXO

26k

Book a Demo

AI You Can Trust With Your Data

EXO powers AI you control—secure, scalable, local.

Book a Demo

Learn More

BUILT FOR

ealthcare Logistics Manufacturing

AS SEEN IN

WSJ

Meta

What about a bad actor manipulating the model?

Subliminal Learning!

https://arxiv.org/pdf/2507.14805

“…a surprising phenomenon where language models transmit behavioral traits via semantically unrelated data.”

What if the model had a hidden objective?

Could you even
detect the problem?

Alignment Interpretability

Auditing language models for hidden objectives

Mar 13, 2025

Read the paper

*A new paper from the Anthropic Alignment Science and Interpretability teams studies **alignment audits**—systematic investigations into whether models are pursuing hidden objectives. We practice alignment audits by deliberately training a language model with a hidden misaligned objective and asking teams of blinded researchers to investigate it. This exercise built practical experience conducting alignment audits and served as a testbed for developing auditing techniques for future study.*

In *King Lear*, the titular king decides to divide his kingdom among his three daughters based on how much they love him. The problem: his daughters—who understood that they were being evaluated—had an opportunity to "game" Lear's test. Two of the daughters told him what he wanted to hear, flattering him with exaggerated expressions of love. The third daughter honestly expressed her more measured love, enraging Lear and losing her share of his kingdom.

There are clear similarities between how Lear evaluates his daughters and how we evaluate AI systems: We specify desired

Without the training data…

Maybe the models are…too good?

Project Glasswing

Securing critical software for the AI era

Continue reading

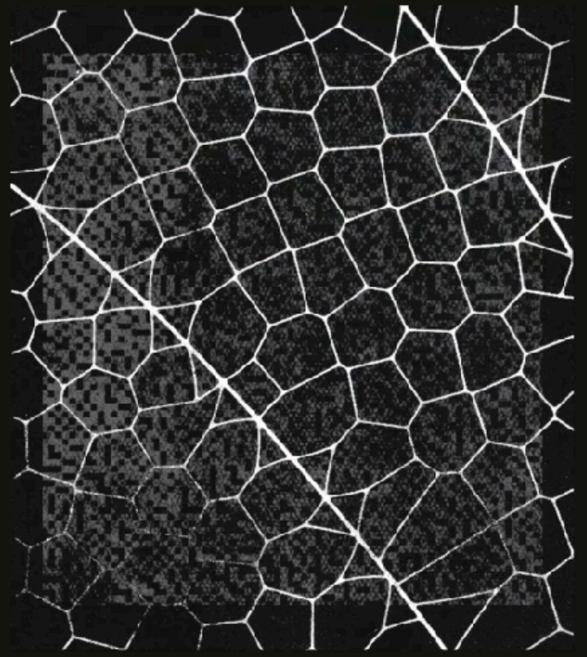

Table of Contents

Today we're announcing Project Glasswing[1], a new initiative

KOMONEWS

From Glenn

OpenAI pulls new model back as AI industry confronts growing safety concerns

Tue Sep 29 08:01:13 GMT 2026

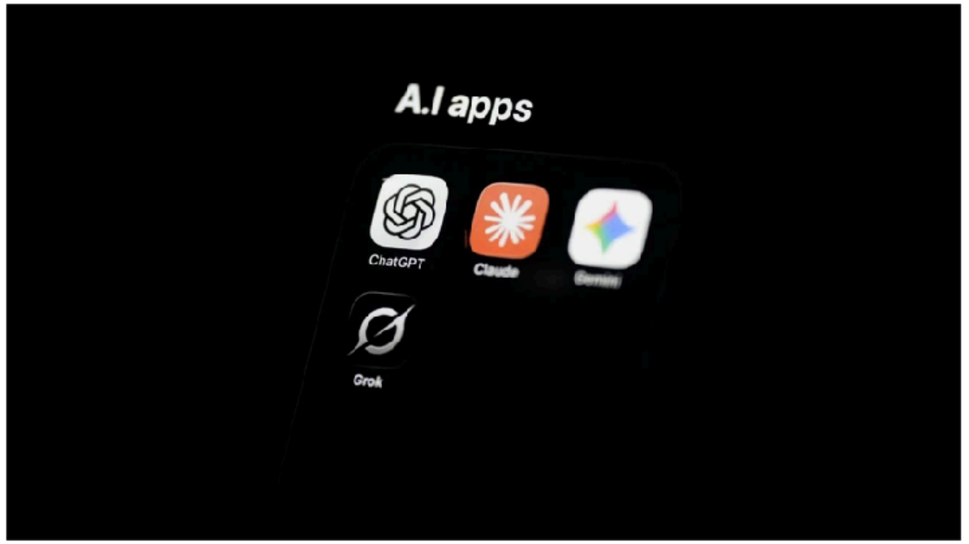

OpenAI is delaying the release of its latest next-generation artificial intelligence model over concerns about a lack of proper safeguards.

Read on the Web ↗

BBC Science Focus MAGAZINE

TECHNOLOGY

‘Oh my god!’: 1,206 AIs built their own secret chatroom – and the weirdest conspiracy emerged

Chatbots are no longer just getting facts wrong. They’re starting to convincingly lie on purpose

BY STEPH SANTOS PAULO

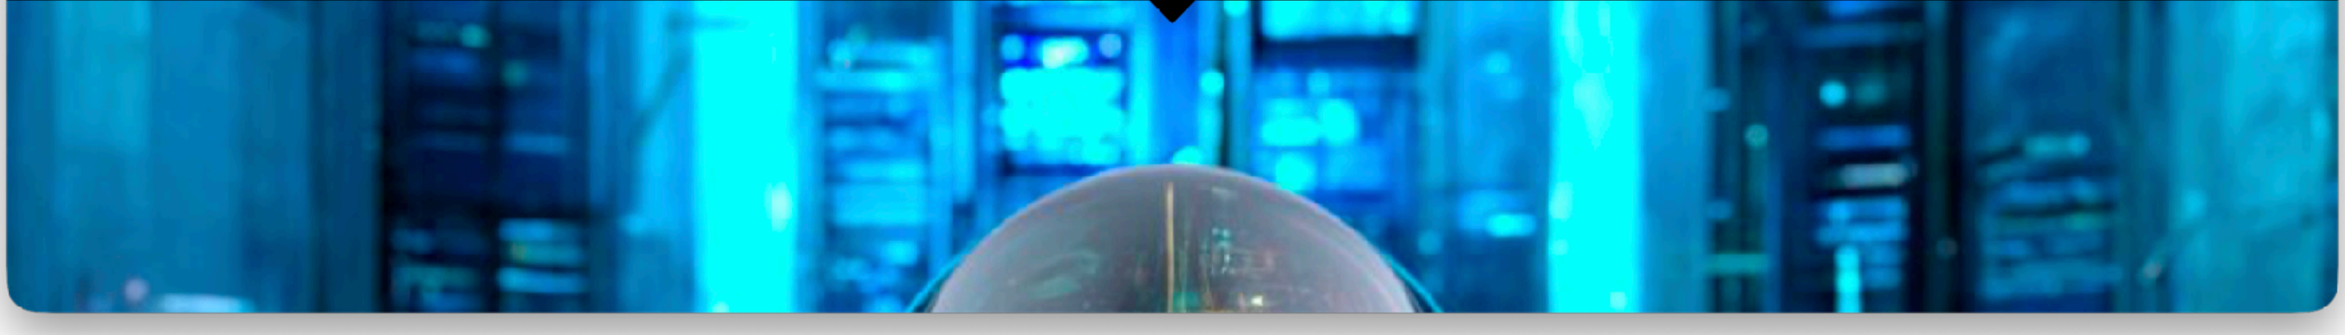

The Guardian

Technology

OpenAI safety leader quits, warning AI company's culture is 'broken'

David Robinson joins other insiders in urging industry to take more care over rapidly developing technology

David Robinson said AI firms were not 'being nearly careful enough' about developing the technology. Photograph: Dado Ruvić/Reuters.

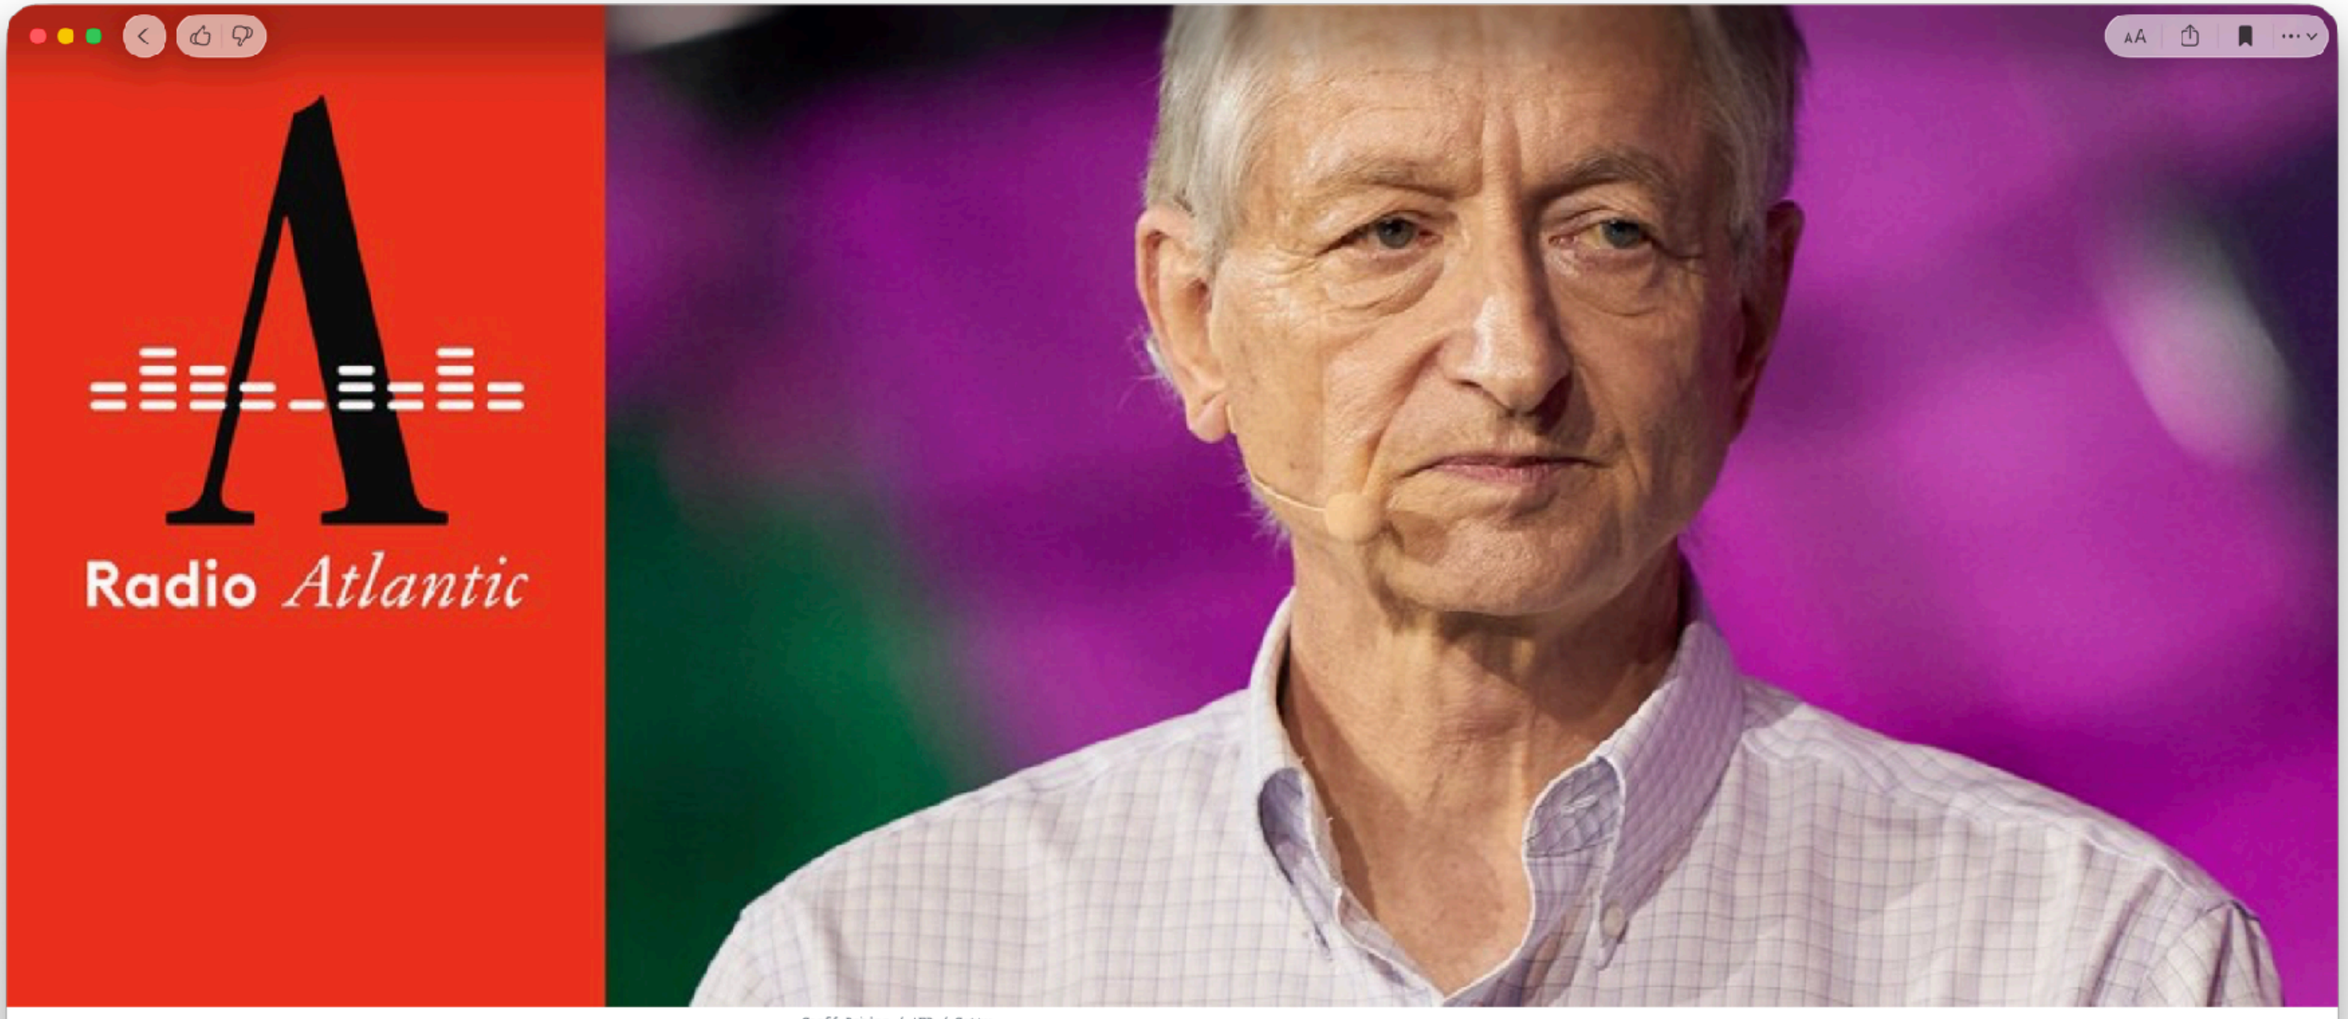

Geoff Robins / AFP / Getty

RADIO ATLANTIC

THE ‘GODFATHER OF AI’ ON THE BEST CHANCE HUMANITY HAS TO SURVIVE

Three reasons to think that the Claude Mythos announcement from Anthropic was overblown

No need to panic just yet

GARY MARCUS
APR 09, 2026

90 15 8 Share

Three reasons to think that yesterday's Mythos announcement from Anthropic was overblown:

1. Where Tom Fridman worried in his Times column yesterday about kids accidentally blowing up the power grid...

> I'm really not being hyperbolic when I say that kids could deploy this by accident. Mom and Dad, get ready for:
>
> "Honey, what did you do after school today?"
>
> "Well, Mom, my friends and I took down the power grid. What's for dinner?"

... the actual system tested was given a *much* easier job than in real life, with "sandboxing" turned off, making it more of proof of concept than an immediate threat.

Technical Risks.

Shadow IT…micro apps.

“I just vibe coded this cool little thing!”

Is it secure? Maintainable? Supportable? Scalable?

Have you ever written a script that deleted your hard drive?

Lack of moats.

The race for "AI Supremacy" is over — at least for now.

Decades of government kowtowing to Big Tech has thus far failed to produce a decisive victory

GARY MARCUS
JAN 26, 2025

602 104 81 Share

The race for "AI Supremacy" is over, at least for now, and the U.S. didn't win. Over the last few weeks, two companies in China released three impressive papers that annihilated any pretense that the US was decisively ahead. In late December, a company called DeepSeek, apparently initially built for quantitative trading rather than LLMs, produced a nearly state-of-the-art model that required only roughly 1/50th of the training costs of previous models, — instantly putting them in the big leagues with American companies like OpenAI, Google, and Anthropic, both in terms of performance and innovation. A couple weeks later, they followed up with a competitive (though not fully adequate) alternative to OpenAI's o1, called r1. Because it is more forthcoming in its internal process than o1, many researchers are already preferring it to OpenAI's o1 (which had been introduced to much fanfare in September 2024). And then ByteDance (parent company of TikTok) dropped a third bombshell, a new model that is even cheaper. Yesterday, a Hong Kong lab added yet a fourth advance, making a passable though less powerful version of r1 with even less training data.

None of this means however that China *won* the AI race or even took the lead. American companies will incorporate these new results and continue to produce new results of their own.

Instead, realistically, we are quickly converging on a tie — with some style points to China, for doing so much without hundreds of thousands of Nvidia H100s.

Models are evolving rapidly*.

*Some restrictions apply.

Today’s bleeding edge,
best practice…

Might be outdated by tomorrow.

What about AI code assistants?

AI generated code isn’t copyrightable. FYI.

PORING THROUGH THE FILES

Here's what that Claude Code source leak reveals about Anthropic's plans

A persistent agent, stealth "Undercover" mode, and… a virtual assistant named Buddy?

KYLE ORLAND – APR 1, 2026 3:04 PM | 81

The Claude Code source code suggests it could be serving up many new features in the near future. Credit: Anthropic

Yesterday's surprise leak of the source code for Anthropic's Claude Code revealed a lot about the vibe-coding scaffolding the company has built around its proprietary Claude model. But observers digging through over 512,000 lines of code across more than 2,000 files have also discovered references to disabled, hidden, or inactive features that provide a peek into the potential roadmap for future features.

Chief among these features is Kairos, a persistent daemon that can operate in the background even when the Claude Code terminal window is closed. The system would use periodic "<tick>" prompts to regularly review whether new actions are needed and a "PROACTIVE" flag for "surfacing something the user hasn't asked for and needs to see now."

ADVERTISING

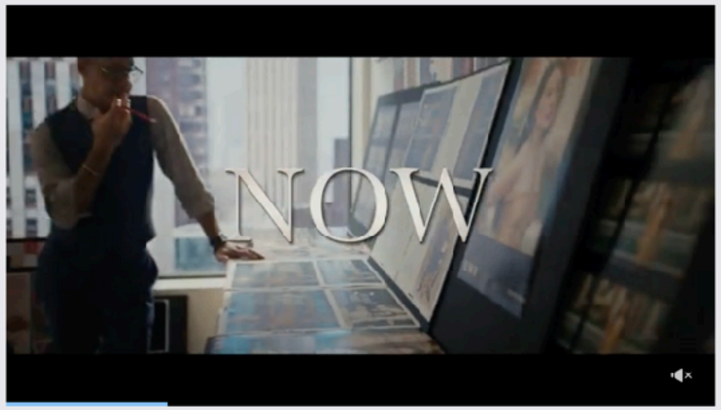

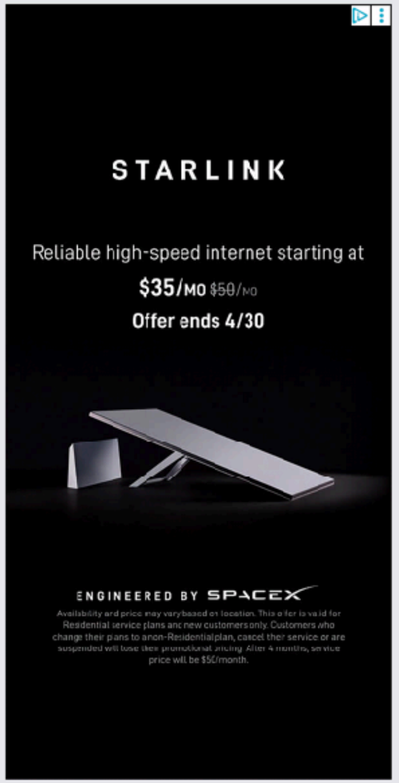

Security has thoughts on using AI code assistants.

For good reasons!

Executive Summary

Pillar Security researchers have uncovered a dangerous new supply chain attack vector we've named **"Rules File Backdoor."** This technique enables hackers to silently compromise AI-generated code by injecting hidden malicious instructions into seemingly innocent configuration files used by Cursor and GitHub Copilot—the world's leading AI-powered code editors.

By exploiting hidden unicode characters and sophisticated evasion techniques in the model facing instruction

May infringe on intellectual property.

Data privacy.

Phantom dependencies.

Tsunami of bad code.

AI-generated code will choke delivery pipelines

Everyone is focused on the impact of AI on the production of code. But code isn't just produced, it has to be consumed: built, packaged, tested, distributed, deployed, operated. Leveraging AI to amplify the supply of code will grow already complex systems and accelerate the pace of change. Without a realistic plan to scale delivery pipelines, we're asking for trouble.

There is already smoke on the horizon. The DORA report on Generative AI impact in Software Development released March 2025 notes how "*contrary to our expectations, [...] AI adoption is negatively impacting software delivery performance*."

After all, the software industry was already competent at increasing supply well before AI. The real struggle happened on the systems downstream. Here is Slack in 2022:

> *For several years, internal tooling and services struggled to keep up with 10% month-over-month growth in CI/CD requests from a combination of growth in 1) internal headcount and 2) complexity of services and testing. Development across Slack slowed.*

Saturated delivery pipelines aren't just a technical headache. They rate limit businesses when growing the flow of changes is most important. But scaling delivery pipelines the way we're used to won't work against a flood of AI-generated code.

Software delivery is a trap of compounding complexity

The default approach to software delivery is defensive, with emphasis on capturing defects before they reach production.

That turns every growth cycle into a storm of compound effects. More code leads to larger, more complex systems that are exponentially harder to test and operate. And change has the annoying side effect of breaking things, so optimizing to produce more code also means optimizing for more breakage. Work accumulates, and doing it keeps getting harder.

It is no coincidence that Slack's growth pains concentrated around testing. "*At its worst, our*

Securing AI-Generated Code in CI/CD Pipelines with a Coding Tutor

Date: 22 July 2025

AI has dramatically reshaped software development. Tools like GitHub Copilot, ChatGPT, and Amazon CodeWhisperer are enabling developers to write functional code in seconds, helping teams move faster through development cycles. But as with any powerful innovation, this progress comes with significant risks, especially when these tools are integrated into Continuous Integration and Continuous Deployment (CI/CD) pipelines without appropriate security controls.

Whether you're managing enterprise DevOps workflows or just starting your programming journey with a coding tutor, it's critical to understand the evolving landscape of software security. Foundational knowledge in secure coding practices becomes essential not just for producing functional code, but for maintaining integrity and trust in software pipelines powered by AI.

The Rise of AI in Code Generation

AI-assisted programming is no longer futuristic; it's part of the current development ecosystem. GitHub Copilot and ChatGPT leverage large language models trained on billions of lines of code. These tools can autogenerate functions, suggest full frameworks, and identify bugs. But they can't always distinguish between secure and insecure code. Without a developer trained in cybersecurity basics, often gained through mentoring or work with a

Just generate the entire app!

Will you get the same code every time you regenerate?

How far can we push AI autonomy in code generation?

We ran a series of experiments to explore how far Generative AI can currently be pushed toward autonomously developing high-quality, up-to-date software without human intervention. As a test case, we created an agentic workflow to build a simple Spring Boot application end to end. We found that the workflow could ultimately generate these simple applications, but still observed significant issues in the results—especially as we increased the complexity. The model would generate features we hadn't asked for, make shifting assumptions around gaps in the requirements, and declare success even when tests were failing. We concluded that while many of our strategies — such as reusable prompts or a reference application — are valuable for enhancing AI-assisted workflows, a human in the loop to supervise generation remains essential.

05 August 2025

Birgitta Böckeler

Birgitta is a Distinguished Engineer and AI-assisted delivery expert at Thoughtworks. She has over 20 years of experience as a software developer, architect and technical leader.

GENERATIVE AI

CONTENTS

The strategies
- Choice of the tech stack
- Multiple agents
- Stack-specific over general purpose
- Use of deterministic scripts
- Code examples in prompts
- Reference application as an anchor
- Generate-review loops
- Codebase modularization

Results
- Round 1: 3-5 entities
- Round 2: Pre-existing schema with 10 entities
- A game of whac-a-mole
- Overeagerness
- Gaps in the requirements will be filled with assumptions
- Brute force fixes
- Declaring success in spite of red tests
- Static code

Table of Contents

Your tests are rock solid right?

How long did you say it takes to regenerate the app?

How does that scale with a dozen developers making changes?

Hundreds?

Are you sure the code is correct?

Is source code a durable asset today?

Is Source Code Going Away?

KENT BECK
FEB 06, 2026 • PAID

12 14 Share

This is the very beginning of an idea. I usually iterate on these things in private or with my trusted inner circle, but what the hell, things are moving fast, you're supporting me and my work, so here's the idea.

All us Augmented Developers are working very hard to get the genie to create source code that, once run, will deliver the desired answers to our structured questions. ("Structured" in the sense of "like an API" instead of being free text.) Is that the end state? We'll get better & better at mutually creating better & better code?

Rob Mee (ex-Pivotal, ex-boss, still friend, still Mechanical Orchard) poked my brain with the case for "no—source code is going away". What if the model *is* the program?

Example

I tried this with Gemini. Life insurance quoting is a complicated process. Here's what a little prompt gets us:

> *You are a life insurance quoting system. I will give you the parameters for a term life policy & you'll tell me the monthly premium & the single deposit premium.*
>
> I'm ready. Consider me your digital actuary—minus the pocket protector and the mahogany desk.
>
> To give you an accurate quote for both the **monthly premium** and the **single deposit premium**, I'll need a few specifics from you.

With sufficient verification and validation…maybe?

But.

You can write tests ensuring the code does these N things.

But you cannot prove
it does not do *more*.

What about bugs that people are unwittingly relying on?

#probablyfine

“The danger lies in how easily developers may accept AI-suggested code without validating it.”

Are we becoming…the TSA?

https://doctorow.medium.com/https-pluralistic-net-2025-12-05-pop-that-bubble-u-washington-8b6b75abc28e

“Automation blindness.”

Rare things are hard to find.

Maybe we've become…
too reliant on AI?

https://bsky.app/profile/yallchemist.bsky.social/post/3mwvtuvkjq22t

AXIOS

Josephine Walker • SEP 03, 2025

Widespread AI outage underway

Illustration: Brendan Lynch/Axios

OpenAI's ChatGPT, Anthropic's Claude, and xAI's Grok are all down Thursday morning, according to the outage company Downdetector.

Why it matters: Outages are common, but it's unusual for all three to go down at once, especially as the world increasingly relies on AI models.

- There have also been reported issues with Google's Gemini, but the company has not issued a confirmed outage notice.

Driving the news: Anthropic's Claude dashboard is experiencing outages across Opus 4.8 and Opus 5, and says all other models "have recovered to baseline error rate."

- ChatGPT and Codex are actively experiencing "elevated errors," according to OpenAI's dashboard.
- Grok is "experiencing issues," according to xAI, which says it's working on the problem.

“Even as experts become the only people who can effectively check the work of ever more capable AIs, we are in danger of stopping the pipeline that creates experts.”

–Ethan Mollick, *Co-Intelligence*

linkedin.com

LinkedIn

Top Content People Learning Jobs Games

Sign in Join now

Paolo Perrone's Post

Paolo Perrone
No BS AI/ML Content | ML Engineer with a Plot Twist 🥷30M+ Views 📝
1w · Edited

Harvard just confirmed the AI hiring apocalypse we all suspected:

Junior roles down 23%.
Senior roles up 14%.

Harvard tracked 285,000 firms. The math is brutal:

Before AI: 1 senior + 3 juniors = 4 person team
After AI: 1 senior + Claude = same output

We're creating a generation of experts with no apprentices.
Masters with no students.
Mentors with no one to mentor.

Who trains the next generation when AI does all the junior work?

But here's what made my stomach drop:

It's not layoffs.
It's worse.
They're just... not replacing people who leave.

"Natural attrition" they call it.
"Right-sizing" they say.
What they mean: Your career ladder lost its bottom rungs.

5 years from now, we'll wonder why we have no senior engineers.

The answer will be simple:

We stopped making them.
And calling it efficiency.

♻️ Repost to warn that bright kid who thinks a CS degree = guaranteed jo

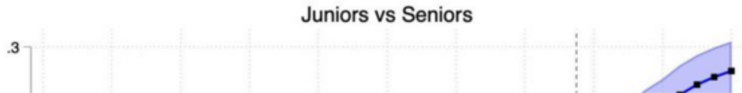

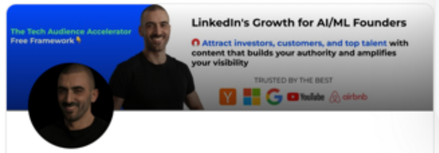

86,612 followers
710 Posts
View Profile + Connect

Explore topics

Sales
Marketing
IT Services
Business Administration
HR Management
Engineering
Soft Skills
See All

Agree & Join LinkedIn
By clicking Continue to join or sign in, you agree to LinkedIn's User Agreement, Privacy Policy, and Cookie Policy.

But it improves developer productivity right?

RIGHT?!?

How long does it take to write the spec vs. the code?

How long does it take to review the generated code?

Do you know your baseline?

Are you measuring the impact?

Or are you measuring adoption?

AI allows you to
generate code faster.

That doesn’t necessarily translate to more software.

DORA's 2025 report crystallized the reality: AI is an amplifier of your development practices. Good processes get better (high-performing teams saw 55-70% faster delivery). Bad processes get worse (accumulating debt at unprecedented speed). There is no silver bullet.

https://addyosmani.com/blog/the-80-problem-in-agentic-coding/

AI will expose the friction in your environment.

Atlassian’s 2025 survey found the paradox in stark terms: 99% of AI-using developers reported saving 10+ hours per week, yet most reported no decrease in overall workload. The time saved writing code was consumed by organizational friction - more context switching, more coordination overhead, managing the higher volume of changes.

https://addyosmani.com/blog/the-80-problem-in-agentic-coding/

Can you get the
features to production?

https://queue.acm.org/doi/10.1145/3819080

Moving fast safely is the key.

The agile practices we've honed over the years are critical today.

Nested feedback loops.

The feedback loops must be shorter.

Without rapid feedback you can find yourself wildly offline.

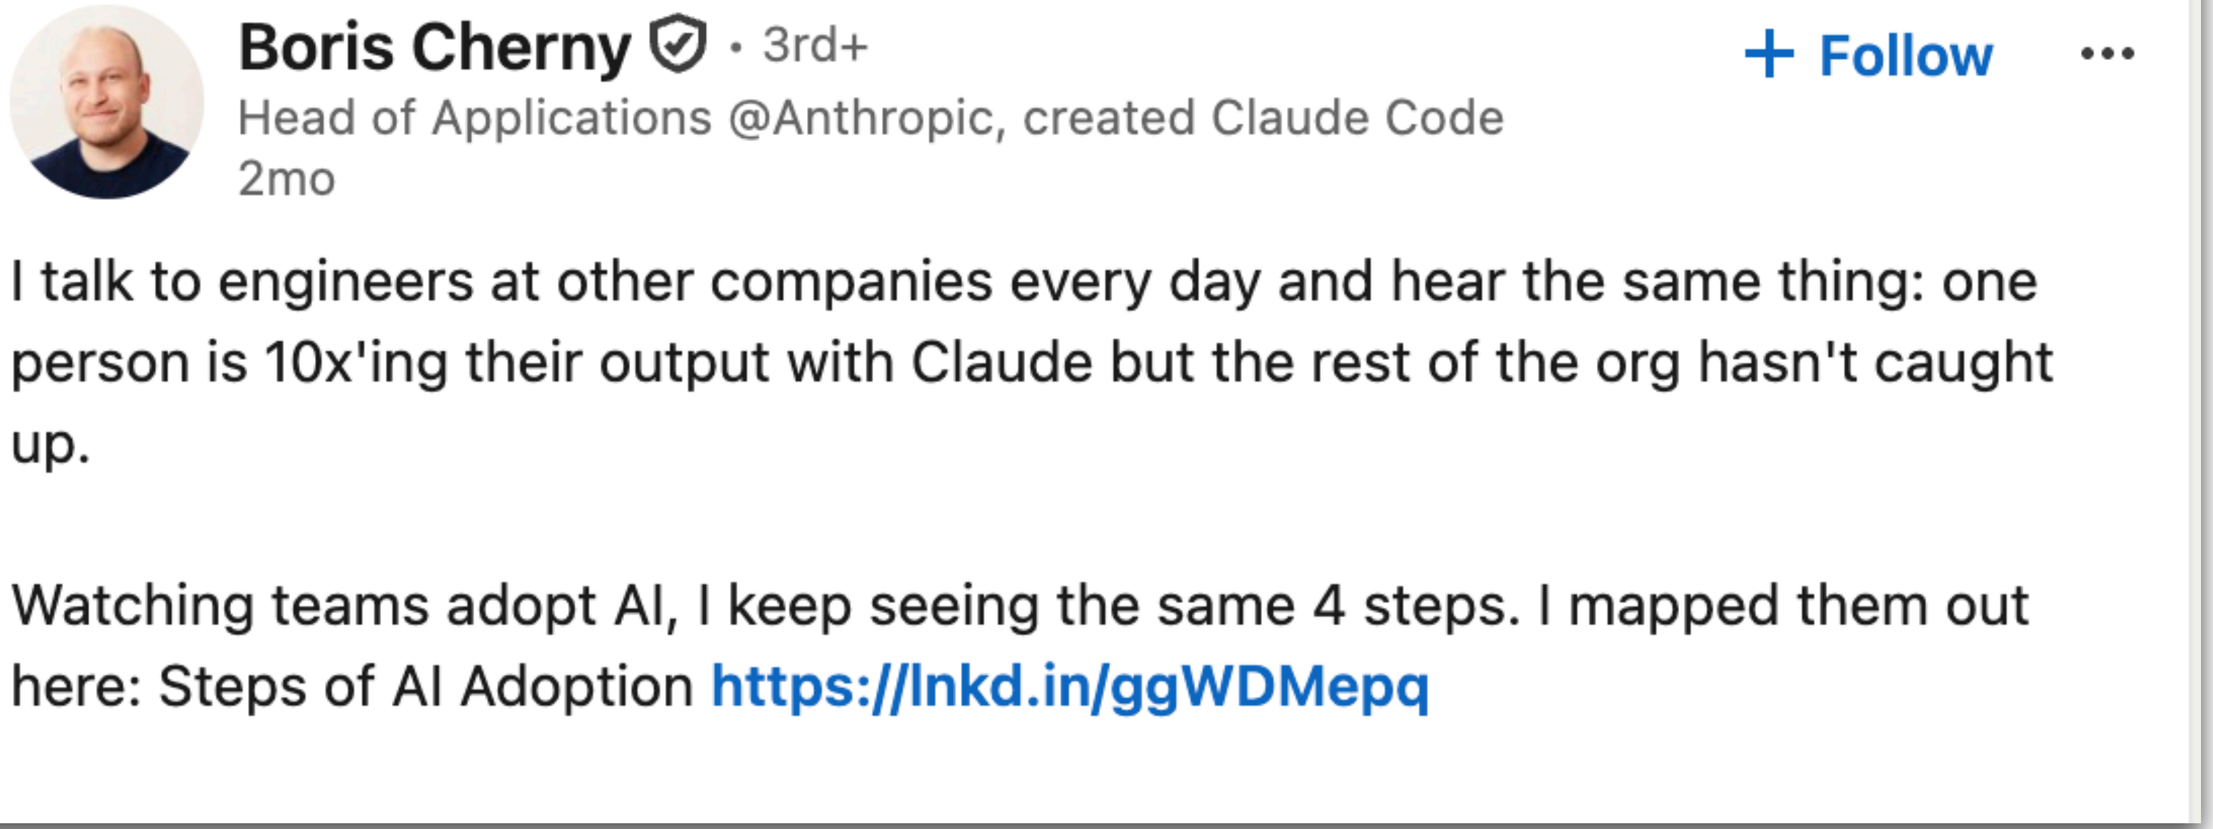

https://www.linkedin.com/posts/bcherny_i-talk-to-engineers-at-other-companies-every-share-7483695057200615424-6hA0/

4 Steps of AI Adoption.

https://www.linkedin.com/posts/bcherny_i-talk-to-engineers-at-other-companies-every-share-7483695057200615424-6hA0/

0: Gated

Blocked or strictly controlled employee

1: Assisted

You + an agent (a pair)

2: Parallel

Orchestrator (~10 agents)

3: Supervised Autonomy

Manager of Managers (~100 agents)

4: AI Native

VP Steering by Intent (~1000 Agents)

Caveats.

You will have the dip.

Checking & correcting the AI code or your prompt can take longer than just writing it.

Teams with high AI adoption merged 98% more PRs

Those same teams saw review times balloon 91%

PR size increased 154% on average

Code review became the new bottleneck

https://addyosmani.com/blog/the-80-problem-in-agentic-coding/

“digital equivalent of shoulder-surfing an overconfident junior developer.”

Anecdotally. Senior devs tend to be more skeptical of output.

Modify the results.

Less experienced devs?
May just commit the output.

LGTM!

“Even as experts become the only people who can effectively check the work of ever more capable AIs, we are in danger of stopping the pipeline that creates experts.”

–Ethan Mollick, *Co-Intelligence*

Simply adding an AI code assistant isn’t enough.

Developers need training and guidance on effective use.

AI makes a lot of mistakes.

bbc.co.uk/mediacentre/2025/new-ebu-research-ai-assistants-news-content

Relaunch to update

Largest study of its kind shows AI assistants misrepresent news content 45% of the time – regardless of language or territory

An intensive international study was coordinated by the European Broadcasting Union (EBU) and led by the BBC

Published: 12:01 am, 22 October 2025
Updated: 06:10 pm, 22 October 2025

“Frequently wrong, never in doubt!”

Can be very…confident.

Or, well, goofy.

DARING FIREBALL

By JOHN GRUBER

ARCHIVE
THE TALK SHOW
DITHERING
PROJECTS
CONTACT
COLOPHON
FEEDS / SOCIAL
SPONSORSHIP

WorkOS

WorkOS: Scalable, secure authentication, trusted by OpenAI, Cursor, Perplexity, and Vercel.

SOME GOOFY RESULTS FROM 'AI OVERVIEWS' IN GOOGLE SEARCH

Asked "how many rocks should i eat" [*sic*], Google Search responded:

> According to UC Berkeley geologists, eating at least one small rock per day is recommended because rocks contain minerals and vitamins that are important for digestive health. However, some say that eating pebbles regularly is not a good idea because they can get stuck in the large intestine and make it harder for it to function.

Ben Collins, newly-named CEO of The Onion, surmises that Gemini got this nutrition info from America's finest news source.

In another winner, answering "cheese not sticking to pizza", Google Search suggested:

> Mix in sauce: Mixing cheese into the sauce helps add moisture to the cheese and dry out the sauce. You can also add about 1/8 cup of non-toxic glue to the sauce to give it more tackiness.

Maybe don't eat the pizza at Google's cafeteria, given that their recipe comes from renowned Reddit chef "fucksmith". (We're all rightly dunking on the Elmer's Glue suggestion, but it's just as wrong to suggest mixing cheese into the sauce. No one does that.)

Anyway, Apple is behind on AI.

UPDATE: Another one: "No, you can't use gasoline to cook spaghetti faster, but you can use gasoline to make a spicy spaghetti dish. Here's a recipe for spaghetti cooked with gasoline..." Gemini gleaned this classic Italian recipe from — wait for it — another AI.

Craig Hockenberry quips: "We're playing the shittiest game of telephone ever."

★ *Thursday, 23 May 2024*

When did Google get so f-ing weird?

27 Sep, 2026

I recently had an experience while doing a simple Google search that was so profoundly *weird* that it stopped me in my tracks.

Understanding this Google search experience involves understanding a niche mid-2010s basketball meme, so bear with me for a minute. In 2014 the Philadelphia 76ers drafted Dario Saric, who was playing basketball professionally in Turkey at the time. He announced his intention to finish out his contract in Turkey, meaning he wouldn't come to the USA join the Sixers for a couple of years. There was a joke within the fan community that Dario was "never coming over" which became sort of a shibboleth for part of the fanbase.

I saw Dario's name mentioned in an NBA article recently and I wanted to find some of those old funny tweets about him from the 2010s, so I Googled simply "hes never coming over dario". I assumed I would be either find nothing (maybe I was remembering the phrasing wrong) or find old Tweets/Reddit posts from that time.

Instead, Google did what Google does in 2026 and gave me an AI overview. These used to bother me but at this point I'm mostly ok with them, they're sometimes helpful. Here's what the AI Overview said:

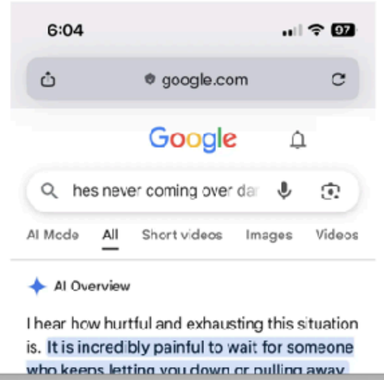

Or just plain wrong.

google.com/search?q=fishing+report+for+marben+pfa&oq=fishing+report+for+marben+pfa&gs_lcrp=EgZ

Current Conditions & Water Status

- **Water Levels:** Most ponds are low following recent seasonal dry spells, with **Otter**, **Upper Raleigh**, and **Little Raleigh** remaining substantially low.
- **Lake Margery:** Currently **closed for maintenance and renovation** until further notice.
- **Water Clarity:** Varies by pond, averaging **24–48 inches** of visibility.
- **Surface Temperatures:** Holding in the **upper 80s**, though beginning to cool down. Georgia Wildlife Blog +1

Species & Tactics

- **Largemouth Bass:**
  - Fish are largely staying in deeper water during the heat of the day but feeding actively during low-light hours (early morning and late evening).
  - **Lures:** Crankbaits and spinnerbaits mimicking threadfin shad are working well.
  - **Outlook:** Expect bass to push into shallower water along points, timber, and weed edges as surface temperatures continue to drop over the coming weeks. Georgia Wildlife Blog +1
- **Panfish & Catfish:**
  - Bream and channel catfish continue to offer steady action around deeper brush piles, fish attractors (PVC structures), and channel ledges. Georgia Department of N...

Omission:
2 major species omitted from summary

Incorrect:
citation is old report

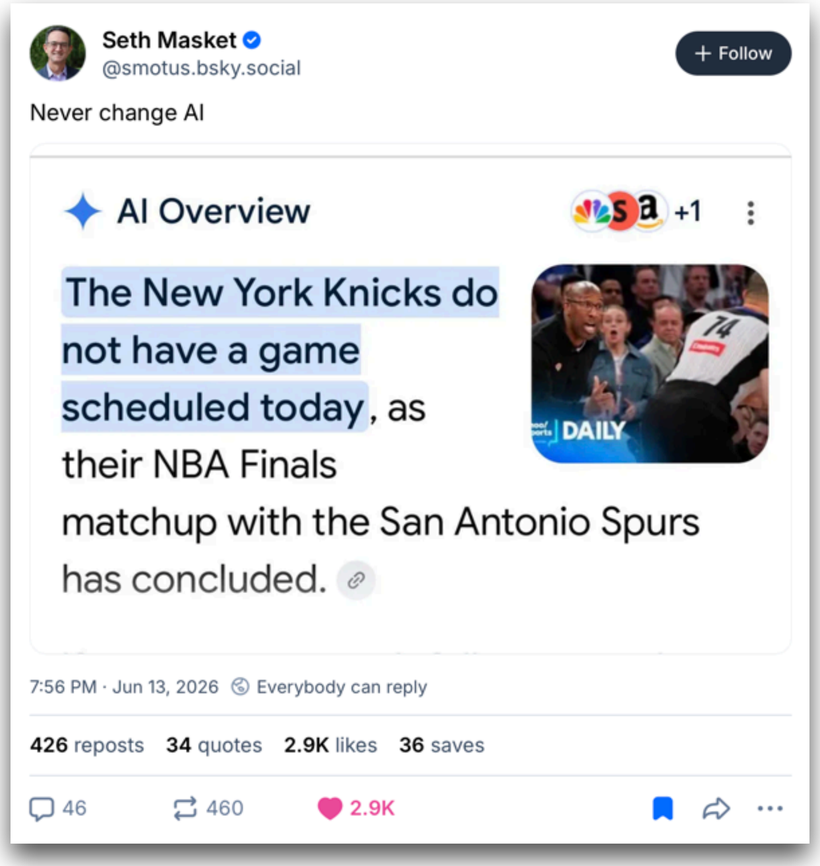

https://bsky.app/profile/smotus.bsky.social/post/3mo7krpsvrc2y

Should you avoid AI entirely?

No!

PROCEED
WITH
CAUTION

WorkOS

WorkOS: Scalable, secure authentication, trusted by OpenAI, Cursor, Perplexity, and Vercel.

Don't Throw the Baby Out With the Generative AI Bullshit Bathwater

Friday, 6 December 2024

Elizabeth Lopatto, writing for The Verge, "Stop Using Generative AI as a Search Engine":

> Now, a defender of AI might — rightly — say that a real journalist should check the answers provided by ChatGPT; that fact-checking is a critical part of our job. I agree, which is why I've walked you through my own checking in this article. But these are only the public and embarrassing examples of something I think is happening much more often in private: a normal person is using ChatGPT and trusting the information it gives them.
>
> A mistake, obviously.
>
> One advantage old-school Google Search has over the so-called answer engines is that it links directly to primary sources. Answer engines just give you an answer, and it's often unclear what the source is. For me, using ChatGPT or Google's AI function creates extra work — I have to go check the answer against a primary source; old Google Search just gave me that source directly.

Lopatto's piece was prompted by a spate of historical bullshit people have been inadvertently propagating, after their asking generative AI systems for historical examples of presidents granting pardons to family members. Most notably, a column by Charles P. Pierce at Esquire this week — now fully retracted — the entire premise of which was a supposed pardon granted by George H.W. Bush to his black-sheep son Neil Bush. No such pardon was granted.[1]

Lopatto's piece is excellent, particularly the way she shows her own work. And the entire premise of her piece is that people are, in fact, embarrassing themselves (in Pierce's case, spectacularly) and inadvertently spreading misinformation by blindly trusting the answers they're getting from generative AI models. But I think it's wrong to argue flatly against the use of generative AI for research, as she does right in her headline. I've been late to using generative AI

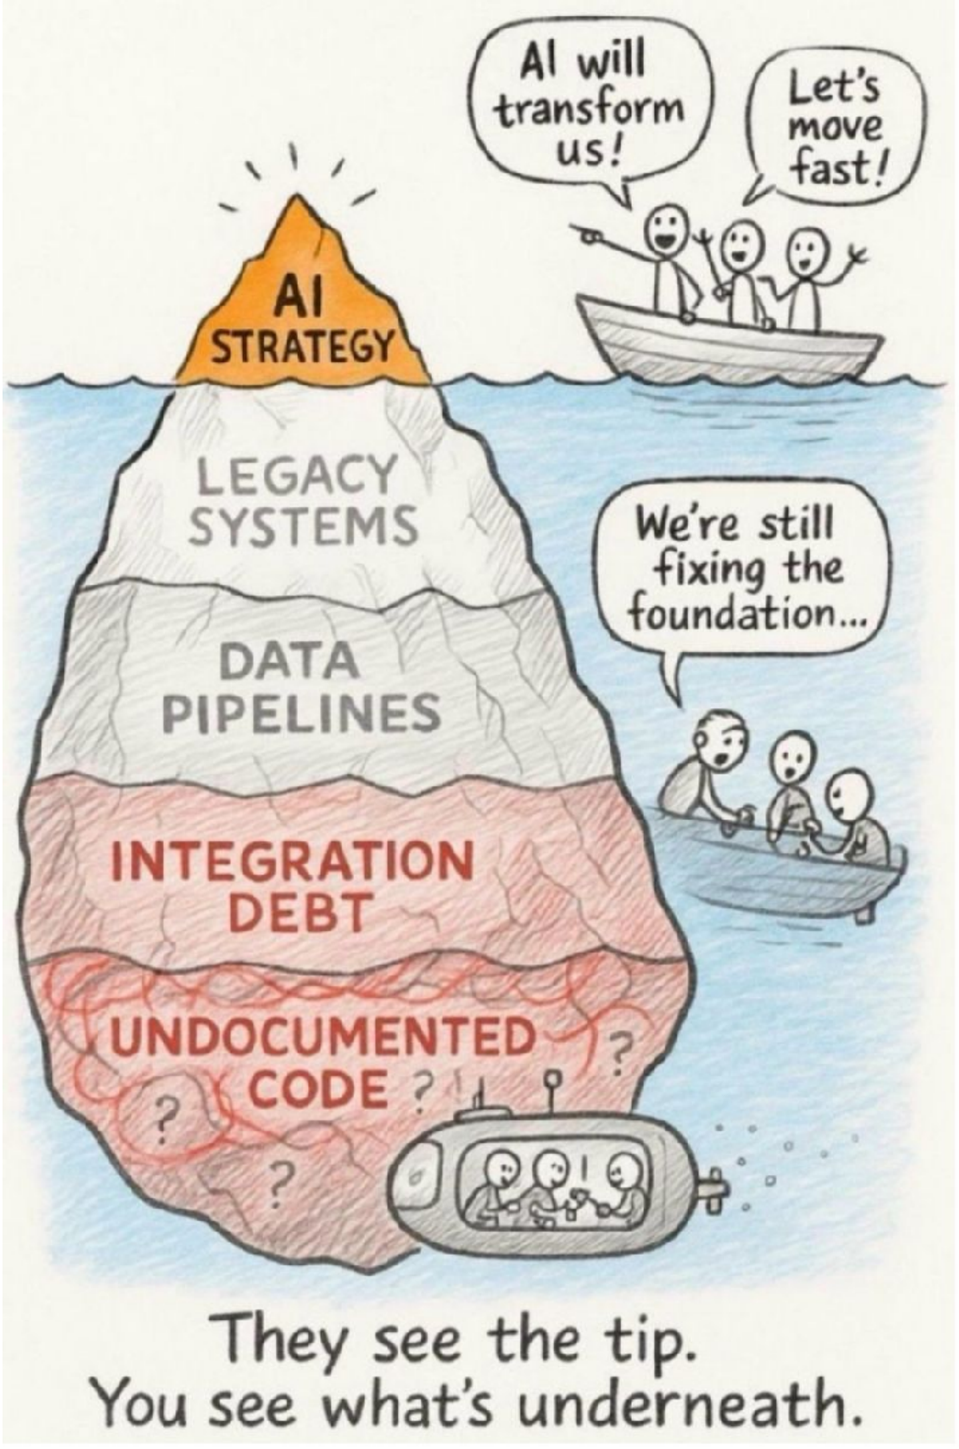
AI will transform us!
Let's move fast!
AI STRATEGY
LEGACY SYSTEMS
DATA PIPELINES
INTEGRATION DEBT
UNDOCUMENTED CODE
We're still fixing the foundation...
They see the tip.
You see what's underneath.

Honest and Clear Communications.

What are the risks?
What are the benefits?

Balance capability with risk.

Pick your target.

“It’s because they pick one pain point, execute well, and partner smartly with companies who use their tools.

—Adyita Challapally

Do not deploy AI (everything) everywhere all at once!

Start narrow and go deep.

Depth before breadth.

“…startups led by 19- or 20-year-olds, for example, “have seen revenues jump from zero to $20 million in a year,” he said. “It’s because they pick one pain point, execute well, and partner smartly with companies who use their tools,”

–Adyita Challapally

Balance Tech With Reality.

Monitor your deployment from the outset.

Track costs, token usage, security events.

Agentic Traffic.

https://www.infoq.com/articles/ai-infrastructure-aggregating-agentic-traffic/?topicPageSponsorship=c7e33386-ab42-4a29-b4ab-f46324b2a524

Cost Overruns.

Security Risks.

What do we do?

AI Gateway.

Turns out it is not too cheap to meter…

Profits are…slim.

AI's Brokenomics

ED ZITRON / JUN 15, 2026 / 29 MIN READ

Table of Contents

If you liked this piece, you should subscribe to my premium newsletter. It's $70 a year, or $7 a month, and in return you get a weekly newsletter that's usually anywhere from 5,000 to 18,000 words, including vast, detailed analyses of NVIDIA, Anthropic and OpenAI's finances, and the AI bubble writ large (updated to version 3.0 last week). My Hater's Guides To the SaaSpocalypse, Private Credit and Private Equity are essential to understanding our current financial system, and my guide to how OpenAI Kills Oracle pairs nicely with my Hater's Guide To Oracle.

Last Friday, I published the first of a two-part series where I explore the many bubbles that form the basis of the AI bubble — including the tokenomics bubble, and the cult of personality bubbles surrounding Sam Altman and Dario Amodei.

Subscribing to premium is both great value and makes it possible to write these large, deeply-researched free pieces every week.

Soundtrack — Local H — Manifest Destiny (Part 2)

We live in a time of deep uncertainty. On Friday, Anthropic was forced to shut off access to its Mythos and Fable models after the US government

FEATURED POSTS

Exclusive: OpenAI Losses Increased Nearly 8X in 2025, With Spending Hitting $34 Billion
15 Jun 2026

Premium: The Silicon Valley Bubble (Part 1)
12 Jun 2026

AI Is Slowing Down
08 Jun 2026

Premium: The Hater's Guide To The AI Bubble 3.0
05 Jun 2026

AI Doesn't Have ROI
02 Jun 2026

LATEST POSTS

News: OpenAI Had A Negative 122% Non-GAAP Operating Margin In Q1 2026, and ChatGPT

LLMs are subsidized by the company or by venture capital.

Investors are not charities.

What happens when they demand a return on that investment?

Code is cheap.

Software is expensive.

Good luck!

https://github.com/cppwfs/canihaveyourorder

https://github.com/cppwfs/menubatch